
Generator-Mediated Bandits: Thompson Sampling for GenAI-Powered Adaptive Interventions

Marc Brooks^{1,*}, Gabriel Durham^{1,*}, Kihyuk Hong^{1,*}, Ambuj Tewari¹

¹Department of Statistics, University of Michigan, Ann Arbor, MI (USA)

{marcbr, gjdurham, kihyukh, tewaria}@umich.edu

Abstract

Recent advances in generative artificial intelligence (GenAI) models have enabled the generation of personalized content that adapts to up-to-date user context. While personalized decision systems are often modeled using bandit formulations, the integration of GenAI introduces new structure into otherwise classical sequential learning problems. In GenAI-powered interventions, the agent selects a query, but the environment experiences a stochastic response drawn from the generative model. Standard bandit methods do not explicitly account for this structure, where actions influence rewards only through stochastic, observed treatments. We introduce *generator-mediated bandit–Thompson sampling* (GAMBITTS), a bandit approach designed for this action/treatment split, using mobile health interventions with large language model-generated text as a motivating case study. GAMBITTS explicitly models both the treatment and reward generation processes, using information in the delivered treatment to accelerate policy learning relative to standard methods. We establish regret bounds for GAMBITTS by decomposing sources of uncertainty in treatment and reward, identifying conditions where it achieves stronger guarantees than standard bandit approaches. In simulation studies, GAMBITTS consistently outperforms conventional algorithms by leveraging observed treatments to more accurately estimate expected rewards.

1 Introduction

Bandit algorithms are ubiquitous in the design of personalized interventions as they provide a flexible framework for learning from context and outcomes over time. Recently, researchers have begun integrating generative artificial intelligence (GenAI) models into these systems to enable richer forms of personalization in areas such as healthcare, e.g., patient care and mobile health (mHealth) [2, 23], personalized education [24], music [48], and marketing [28]. By generating rich, tailored content on-the-fly, GenAI integration expands the scope of personalization beyond what fixed libraries can offer, since predefined intervention options are unable to accommodate the full range of user contexts encountered in practice. While this real-time generative capability opens new possibilities for intervention design, it disrupts the action/reward feedback loop that underpins traditional bandit learning. Rather than directly acting on the environment, the agent selects a query—e.g., a prompt to a large language model (LLM)—and the environment receives the stochastic output generated in response to that query, which we refer to as the “treatment.” This separation between the agent’s action and the delivered treatment introduces additional stochasticity into the intervention process and motivates the need for algorithms that explicitly account for the generative mechanism.

Broadly, we focus on online decision-making problems where the treatment delivered to the environment is a random function of the agent’s action. As a motivating case study, we consider the use of LLMs in sequential mHealth interventions. Traditionally, mHealth systems deliver text-based

*Equal contribution.

support using data collected from smartphones and wearables, adapting to rapid fluctuations in user context, physiology, and behavior [39, 50]. Bandit algorithms have become a workhorse for this sequential personalization, with personalized Just-in-Time Adaptive Interventions (pJITAIs) serving as a popular setting in which such methods are developed and evaluated [51, 33, 14].

This framework for mHealth intervention has inspired countless recent advances in bandit algorithms; e.g., allowing partial pooling of information across users [19, 52, 20], accommodating non-stationarity to adapt to shifting user behavior patterns [33], and integrating high-dimensional covariates [3]. While these advancements further optimize bandit learning within this intervention framework, such approaches are fundamentally limited by reliance on a fixed library of text intervention options.

Dynamic message generation via LLMs offers a path forward. Throughout our analyses, we assume access to a generator whose responses are guaranteed to be safe and suitable for delivery (see Section 7). In this setting, the agent selects a query based on user context, submits both query and user context to the LLM, and delivers the generated text response as the treatment. The user then produces a reward signal (e.g., behavior or engagement). Crucially, while the agent selects the query, it cannot control the LLM’s response.

This setup motivates a bandit framework where the agent’s action affects the environment only through a high-dimensional, stochastic treatment. Standard bandit algorithms are not designed to exploit this indirect treatment mechanism, leading to inefficient learning by failing to make use of the structure linking treatments to rewards. Section 3 more precisely defines this *generator-mediated bandit* (GAMBIT) framework.

We develop a Thompson sampling-based framework tailored to this action/treatment split, incorporating explicit models of the treatment and reward generation processes. This framework serves as the backbone of our approach due to its central role in pJITAI designs, which often rely on Thompson sampling both for its theoretical guarantees and for its compatibility with downstream causal analyses using logged data, a key consideration in mobile health research [33]. We call the proposed approach *generator-mediated bandit–Thompson sampling* (GAMBITTS).

GAMBITTS selects actions by reasoning over both the distribution of treatments induced by each query and the reward those treatments generate, using a two-stage sampling procedure that captures uncertainty across the intervention process. In settings where the agent has simulation access to the treatment-generating mechanism, this structure enables offline estimation of the treatment model, accelerating online adaptation. To address the high dimensionality of treatments in generative settings (e.g., LLM-generated text), GAMBITTS projects each observed treatment into a fixed low-dimensional representation, enabling scalable and sample-efficient learning.

This work responds to the core challenge of learning when actions yield stochastic, observed treatments. The primary contributions are: (a) the GAMBITTS framework for bandit learning with stochastic treatments, (b) algorithmic instantiations of GAMBITTS that vary in their approach to modeling the treatment-generation process, ranging from fully online updates to approaches that pretrain the treatment model, as well as ensemble-based variants that support nonlinear reward modeling, (c) theoretical regret bounds for GAMBITTS, including decompositions that isolate contributions from treatment and reward uncertainty, showing when modeling treatment generation leads to improved learning efficiency, and extending guarantees to flexible, nonlinear reward models, (d) empirical results demonstrating the limitations of standard methods and the performance gains achieved by GAMBITTS-based approaches.

2 Related Work

In the bandit literature, researchers have illustrated how leveraging the causal structure underlying sequential interventions can improve policy learning [30, 37, 35]. Subsequent work has supported this conclusion, developing bandit algorithms in a wide variety of causal structures [57, 56]. These approaches model a set of binary, interconnected random variables $\mathbf{X} = \{X_1, \dots, X_N\}$ with general causal pathways to a reward Y . Actions are framed as interventions on subsets of \mathbf{X} , and algorithms are aimed at optimizing reward in settings where the action space corresponds to modifying causally linked features that influence Y . In parallel, Sen et al. [2017] and Maiti et al. [2022] explore related environments where unobserved confounding can impact bandit learning [45, 38].

While the works discussed above emphasize general causal structures with limited treatment representations, our focus is the inverse: a specific mediation structure, where actions influence reward through an observed stochastic intermediate, but with rich, high-dimensional treatments (i.e., the intermediate generative output). This designed mediation, where the effect of the agent’s action operates through the generator’s stochastic response, parallels the instrumental variable (IV) framework, in which an instrument affects outcomes only indirectly via a stochastic channel [22]. Recent work has connected IV estimation and bandit design in classical econometric settings [58, 9], while related research has explored similar structures under noncompliance, where the realized action may differ from the one selected by the agent [49, 29]. Kallus [2018] makes this connection explicit through the *instrument-armed bandit* framework [25]. Most recently, Zou et al. [2025] study mediated bandit environments from a psychometric perspective [59].

Appendix A discusses further connections and distinctions between our framework and the methods discussed above, along with recent work on LLMs and bandits and related, application-focused, literature from the mHealth community.

3 Problem Formulation and Notation

We consider a sequential decision-making setting in which, at each time $t = 1, \dots, T$, an agent: (i) observes a context $X_t \in \mathcal{X} \subseteq \mathbb{R}^{d_C}$, (ii) selects an action $A_t \in \mathcal{A}$, where $|\mathcal{A}| = K \in \mathbb{N}$, (iii) sends (A_t, X_t) to a generator, (iv) receives response $G_t \in \mathcal{G}$ and delivers G_t to the environment, and (v) observes reward $Y_t \in \mathbb{R}$. The goal of the agent is to learn a policy for choosing actions that maximizes cumulative expected reward over time.

As a concrete example, consider an mHealth intervention designed to encourage physical activity. At time t , the agent: (i) observes user context/covariates X_t (e.g., steps taken the previous day and current location), (ii) selects an action A_t , corresponding to a query from a finite set of options (e.g., “Please write a message that encourages the user to walk more; I have attached the user’s current location and previous day’s step count”), (iii) sends the query, along with the context, to an LLM, (iv) delivers the LLM’s response, G_t , to the user, and (v) observes reward Y_t (e.g., user’s subsequent step count).

In our motivating example, G represents the output of an LLM, and the treatment space \mathcal{G} consists of natural language text. In such GenAI settings, \mathcal{G} is exceedingly high-dimensional, making it natural to posit the effect of G on Y occurs through a lower-dimensional embedding $Z^* \in \mathcal{Z}^* \subseteq \mathbb{R}^{d^*}$, which serves as a sufficient representation of the treatment for predicting reward. Recent work in high-dimensional causal inference supports this assumption, showing that such lower-dimensional embeddings can yield representations sufficient for treatment effect estimation (see, e.g., Veitch et al. [2019]) [53]. Because Z^* is not observed in practice, the analyst must instead specify a working embedding $h(G) = Z \in \mathcal{Z} \subseteq \mathbb{R}^d$. We discuss strategies for constructing h in Section 4. Furthermore, in practice, the agent may also restrict which components of X_t are passed to the generator; we discuss this case in Appendix B.1.

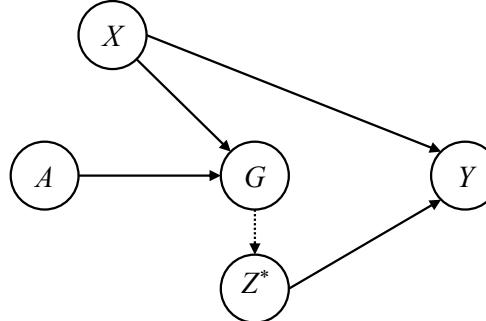


Figure 1: Generator-Mediated Bandit Causal Structure

Notes: (i) Dotted lines represent deterministic relationships.

(ii) Arrows from X to A are omitted to represent the data-generating process rather than the decision logic.

(iii) A : Action/Query; X : Full Context; G : Generated Response; Z^* : True Response Embedding; Y : Reward.

We refer to the resulting learning setup as a *generator-mediated bandit* (GAMBIT) problem. This causal structure is summarized in the directed acyclic graph (DAG) in Figure 1.

4 Methodology: Generator-Mediated Bandit–Thompson Sampling

As discussed above, we seek a Thompson sampling-based framework for our bandit agent, in line with its prevalence in mHealth interventions. We observe two key conditional structures to model:

$$Z \mid A, X, \text{ and } \mathbb{E}[Y \mid Z, X],$$

which correspond to the treatment and reward models. Modeling these two components comprises the two stages of the proposed Thompson sampler and induces the following parametrized environment:

$$\begin{array}{ll} \text{Treatment model: } Z_t \sim f_1(z; A_t, X_t, \theta_1) & \text{Treatment prior: } \theta_1 \sim \pi_1, \\ \text{Reward model: } \mathbb{E}[Y_t \mid Z_t, X_t] = m_2(Z_t, X_t; \theta_2) & \text{Reward prior: } \theta_2 \sim \pi_2. \end{array}$$

The construction of working treatment representation $Z = h(G)$ depends on the application domain. In marketing applications, for instance, G may represent AI-generated images, while in our motivating mHealth setting, G corresponds to LLM-generated text, necessitating different approaches to representation. For implementation, all algorithmic steps operate on the observed working embedding Z , not the unobserved benchmark Z^* . While the GAMBITTS framework is agnostic to the form of h , the choice can substantially influence empirical performance. The effectiveness of the GAMBITTS projection approach depends on how well Z preserves reward-relevant information; when $Y \perp\!\!\!\perp Z^* \mid Z, X$, modeling $Y \mid Z, X$ captures the full signal in Z^* (and thus that in G as well). As the conditional dependence $Y \mid Z, X$ diverges from $Y \mid Z^*, X$, performance may degrade. We explore this phenomenon empirically in Section 6 and further discuss robustness to embedding misspecification in Appendix C. In our motivating setting, many projection strategies have been proposed for natural language-based interventions; see, e.g., [18, 53, 43, 11, 21, 54].

We propose two general approaches for GAMBITTS implementation: a *fully online* variant that updates both treatment and reward models sequentially, and a *partially online* variant that pretrains the treatment model using offline data and updates only the reward model during deployment. To improve flexibility and accommodate nonlinear reward relationships, we extend both variants with the ensemble-based sampling strategies introduced in Lu and Van Roy [2017] [36].

The algorithms below focus on settings where the decision to intervene (e.g., to send the user a text message) has already been made. In practice, mHealth interventions often include an initial decision of whether to send a message at all; GAMBITTS can be trivially extended to accommodate this preliminary selection step, as discussed in Appendix B.2.

4.1 Fully Online Generator-Mediated Bandit–Thompson Sampling

Let $\mathcal{D}_t = \{X_s, A_s, G_s, Y_s\}_{s=1}^{t-1}$ represent the data collected up to time t . Algorithm 1 below describes a fully online approach for selecting action A_t based on \mathcal{D}_t .

Algorithm 1 Fully Online GAMBITTS (foGAMBITTS)

- Inputs:** Data \mathcal{D}_t , priors π_1, π_2 , models f_1, m_2 , current context x_t .
1: **Derive** posterior distributions $P_1(\theta_1 \mid \mathcal{D}_t)$ and $P_2(\theta_2 \mid \mathcal{D}_t)$.
2: **Sample** $\theta_{1,t} \sim P_1(\theta_1 \mid \mathcal{D}_t)$ and $\theta_{2,t} \sim P_2(\theta_2 \mid \mathcal{D}_t)$.
3: **For** each $a \in \mathcal{A}$, calculate

$$\mathbb{E}_{\theta_{1,t}, \theta_{2,t}}[Y \mid a, x_t] = \int_Z m_2(z, x_t; \theta_{2,t}) f_1(z; a, x_t, \theta_{1,t}) dz.$$

- 4:** **Set** $A_t = \arg \max_{a' \in \mathcal{A}} \mathbb{E}_{\theta_{1,t}, \theta_{2,t}}[Y \mid a', x_t]$.

The integral in Step 3 can often be intractable in closed form, but can be efficiently approximated via Monte Carlo approaches. Furthermore, because GAMBITTS selects agents via posterior sampling, it

induces well-defined randomization probabilities that can be used for downstream causal inference or offline evaluation. Specifically, the probability of selecting action a at time t is

$$p_{t,a}^* = \int_{\Theta_2} \int_{\Theta_1} \mathbb{1}_{\left\{ a = \arg \max_{a' \in \mathcal{A}} \mathbb{E}_{\theta_1, \theta_2} [Y | a', x_t] \right\}} P_1(\theta_1 | \mathcal{D}_t) P_2(\theta_2 | \mathcal{D}_t) d\theta_1 d\theta_2.$$

In settings where these probabilities are required for post hoc causal analyses (e.g., analysis of pJITAIs), the full set $\{p_{t,a}^*\}_{a \in \mathcal{A}, t=0, \dots, T}$ can be approximated retrospectively, after the study period, when additional computational resources are available [4, 33]. Alternatively, if sufficient computing power is available at deployment time, the agent may compute these probabilities online and sample from them directly. As with the integral in Step 3, the quantities $p_{t,a}^*$ can be estimated via Monte Carlo sampling from the posterior distributions over θ_1 and θ_2 .

4.2 Partially Online Generator-Mediated Bandit-Thompson Sampling

In settings where the agent has simulation access to the generator, or access to additional data sources, it may be possible to learn the treatment model independently of user interactions. For example, in our motivating application, an mHealth intervention designer could query an LLM multiple times without sending any generated text to users.

This capability motivates a *partially online GAMBITTS* (poGAMBITTS). The poGAMBITTS framework exploits simulation access to the generator to offload learning of the first-stage distribution, $Z | A, X$ to an offline phase. Empirically, this decomposition can substantially improve online performance, as discussed in Section 6.

Here, we consider the generator G to be a random function, which the agent can query using prompt-context pairs $(a, x) \in \mathcal{A} \times \mathcal{X}$. In the poGAMBITTS framework, the agent can obtain an offline approximation of the treatment distribution f_1 by: (i) simulating M_{off} prompt-context pairs $\{(a^1, x^1), \dots, (a^{M_{off}}, x^{M_{off}})\}$, (ii) using the generator to simulate $G(a^i, x^i)$, with $g^i(a^i, x^i)$ denoting the observed output and z^i denoting its observed \mathcal{Z} embedding, for $i = 1, \dots, M_{off}$, and (iii) taking $f_1^{off}(z; A_t, X_t)$ to be the empirical conditional distribution observed in $\mathcal{D}^{off} := \{(a^i, x^i, z^i)\}_{i=1}^{M_{off}}$.

In cases where $|\mathcal{X}|$ is finite, the agent can repeatedly prompt the generator for each prompt-context pair $(a, x) \in \mathcal{A} \times \mathcal{X}$ and directly use the empirical distribution of $Z | A, X$ in \mathcal{D}^{off} as f_1^{off} . When $|\mathcal{X}|$ is infinite, the agent must instead fit a model to estimate the conditional distribution of $Z | A, X$ based on \mathcal{D}^{off} , and use this estimate as the treatment model.

Algorithm 2 Partially Online GAMBITTS (poGAMBITTS)

- Inputs:** Data \mathcal{D}_t , prior π_2 , model m_2 , current context x_t , pretrained model f_1^{off} .
- 1: **Derive** the posterior distribution $P_2(\theta_2 | \mathcal{D}_t)$.
 - 2: **Sample** $\theta_{2,t} \sim P_2(\theta_2 | \mathcal{D}_t)$.
 - 3: **For** each $a \in \mathcal{A}$, calculate

$$\mathbb{E}_{\theta_{2,t}} [Y | a, x_t] = \int_{\mathcal{Z}} m_2(z, x_t; \theta_{2,t}) f_1^{off}(z; a, x_t) dz.$$

- 4: **Set** $A_t = \arg \max_{a' \in \mathcal{A}} \mathbb{E}_{\theta_{2,t}} [Y | a', x_t]$.

As in foGAMBITTS, the probability of selecting action a at time t can be written explicitly, here as

$$\int_{\Theta_2} \mathbb{1}_{\left\{ a = \arg \max_{a' \in \mathcal{A}} \mathbb{E}_{\theta_2} [Y | a', x_t] \right\}} P_2(\theta_2 | \mathcal{D}_t) d\theta_2,$$

which can be computed or approximated retrospectively if randomization probabilities are needed for post hoc causal analyses. Again, we can approximate $\mathbb{E}_{\theta_{2,t}} [Y | A_t = a, X_t = x_t]$ via Monte Carlo.

Moreover, one may wish to blend the approaches in Algorithms 1 and 2 by using simulated offline data to form a posterior distribution over the treatment model parameters, $P(\theta_1 | \mathcal{D}^{off})$. During online deployment, the agent may either continue updating this posterior online, or fix it and sample throughout learning (avoiding the computational cost of full online θ_1 posterior updates).

4.3 Ensemble-Based GAMBITTS Approaches

Algorithms 1 and 2 require explicit posterior updates for reward model parameters (θ_2), restricting the class of models that can be feasibly used. This restriction can be particularly limiting for two reasons. First, flexible models are often desirable to avoid strong parametric assumptions on $\mathbb{E}[Y | Z, X]$. Second, when Z is an imperfect proxy for Z^* , flexible models may help capture residual information and improve the approximation of $\mathbb{E}[Y | Z^*, X]$ through $m_2(Z, X)$.

To address these challenges, we propose an ensemble-based posterior approximation strategy, drawing on the ensemble sampling framework introduced in Lu and Van Roy [2017] [36]. This approach maintains an ensemble of models to approximate the posterior distribution over parameters, making posterior sampling tractable even for complex or nonlinear models. We present ensemble-based variants of Algorithms 1 and 2 (ens-foGAMBITTS and ens-poGAMBITTS, respectively) in Appendix D. In Section 6, we evaluate the performance of ens-poGAMBITTS using a multilayer perceptron (MLP) model for $\mathbb{E}[Y | Z, X]$.

5 Theoretical Guarantees

This section explores regret guarantees for the GAMBITTS algorithms introduced in Section 4. We revisit the generator-mediated bandit learning problem from Section 3, now from a theoretical perspective. Throughout, we assume a finite context space (i.e., $|\mathcal{X}| = C < \infty$) and focus on the correctly specified setting where the working treatment projection matches the true underlying treatment embedding (i.e., $Z = Z^*$).

Under these assumptions, treatment and reward generation at time t follows a two-stage stochastic process. First, given context x_t and action a_t , the generator samples an intermediate treatment variable $z_t \in \mathbb{R}^d$ from a distribution $\xi_{x_t, a_t}^{\theta_1}$ (with density $f_1(z | x_t, a_t, \theta_1)$) supported on the ball $\mathbb{B}_d(B) = \{z \in \mathbb{R}^d : \|z\|_2 \leq B\}$. Second, the agent receives a noisy reward $y_t = m_2(z_t, x_t; \theta_2) + \eta_t$, where η_t is σ_2 -subgaussian noise. Here, $\theta = \{\theta_1, \theta_2\}$ are unknown parameters. We also assume that the mean reward distribution $m_2(Z, x; \theta_2)$, under $Z \sim \xi_{x, a}^{\theta_1}$, is σ_1 -subgaussian for all $(x, a) \in \mathcal{X} \times \mathcal{A}$.

The goal of the agent is to minimize the time- T cumulative Bayesian regret, defined as

$$\text{BR}_T := \mathbb{E} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{x_t, A_t^*}^{\theta_1}} [m_2(Z, X_t; \theta_2) | X_t, A_t^*, \theta] - \mathbb{E}_{\xi_{x_t, A_t}^{\theta_1}} [m_2(Z, X_t; \theta_2) | X_t, A_t, \theta] \right) \right],$$

where the inner expectations are over Z and the outer expectations over X , θ_1 , θ_2 , and agent actions A . The benchmark a_t^* denotes the Bayes-optimal action in context x_t , defined by $A_t^* := \arg \max_{a \in \mathcal{A}} \mathbb{E}_{\xi_{x_t, a}^{\theta_1}} [m_2(Z, x_t; \theta_2) | \theta]$.

As a baseline, we consider a Thompson sampling agent that models the problem as a conventional multi-armed contextual bandit over $\mathcal{X} \times \mathcal{A}$, ignoring the generator-mediation structure. We refer to this as a “standard” Thompson sampling approach, as it captures a natural modeling choice that treats rewards as a direct consequence of actions. This serves as a point of comparison in both our theoretical and empirical results. As shown in Appendix E, its Bayesian regret satisfies

$$\text{BR}_T^{standard} \leq \tilde{\mathcal{O}}((\sigma_1 + \sigma_2) \sqrt{CKT}) \quad (1)$$

While the regret bound for the standard Thompson sampling agent is tight (i.e., matching known lower bounds), it becomes unfavorable when either the context or action space is large. This motivates the GAMBITTS approach, which can exploit shared structure in the treatment space to support generalization across arms. We begin by analyzing poGAMBITTS, as this approach aligns more closely with our motivating setting, then turn to the foGAMBITTS variant. Appendix E presents the proofs of all results discussed below, along with additional theoretical results.

5.1 Regret Analysis for poGAMBITTS

We begin by analyzing the regret of poGAMBITTS. This algorithm uses empirical conditional distributions, with densities $f_1^{off}(z; x, a)$, that approximate the true treatment distributions $\xi_{x,a}^{\theta_1}$ and define the estimated mean reward

$$\hat{m}_2(x, a; \theta_2) = \int_{\mathcal{Z}} m_2(z, x; \theta_2) f_1^{off}(z; x, a) dz.$$

poGAMBITTS can be viewed as a Thompson sampling algorithm operating on the estimated mean reward function \hat{m}_2 . More generally, we can consider other traditional bandit algorithms operating in the ‘‘partially online stochastic treatment’’ setting, where actions are selected based on an estimated reward model derived from an empirical treatment distribution. Although actions are chosen with respect to \hat{m}_2 , rewards are generated according to the true mean m_2 , introducing a form of model misspecification. Nevertheless, we show that if the empirical treatment model is sufficiently close to the true distribution, then running a no-regret contextual bandit algorithm on \hat{m}_2 still yields no-regret guarantees for the original problem.

In the following results, we let $\varepsilon > 0$ and consider empirical treatment models f_1^{off} satisfying

$$KL(f_1^{off}(\cdot; x, a) \| f_1(\cdot; x, a, \theta_1)) \leq \varepsilon \quad \text{for all } (x, a) \in \mathcal{X} \times \mathcal{A}. \quad (2)$$

As shown in Appendix E, under regularity conditions on $\xi_{x,a}^{\theta_1}$, f_1^{off} will satisfy (2) given $\text{poly}(d, \varepsilon^{-1})$ draws from the simulator for each $(x, a) \in \mathcal{X} \times \mathcal{A}$.

Theorem 1. *Let f_1^{off} be an empirical model for the treatment distribution satisfying (2). Let Alg be a contextual bandit algorithm that selects actions based on the estimated reward function \hat{m}_2 . Then, running Alg in the partially online stochastic treatment setting gives T-step Bayesian regret BR_T^{Alg} , with*

$$BR_T^{\text{Alg}} \leq \mathcal{O}(\widehat{BR}_T^{\text{Alg}} + T\sqrt{\varepsilon T}),$$

where $\widehat{BR}_T^{\text{Alg}}$ is the T-step Bayesian regret of Alg with respect to \hat{m}_2 .

Theorem 1 bounds the agent’s true regret in terms of its regret under the estimated reward model induced by its misspecified treatment distribution. For poGAMBITTS with a linear reward model, we get the following result.

Corollary 1. *Let f_1^{off} satisfy (2). If $m_2(z, x_t; \theta)$ is linear in $z \in \mathbb{R}^d$, then the Bayesian regret of poGAMBITTS (Algorithm 2) satisfies $BR_T^{\text{po:lin}} \leq \tilde{\mathcal{O}}(\sigma_2 d \sqrt{T} + T \sqrt{\varepsilon T})$.*

When $\varepsilon \leq (\sigma_2 d / T)^2$, the regret simplifies to $\tilde{\mathcal{O}}(\sigma_2 d \sqrt{T})$. This is sharper than the standard Thompson sampling bound whenever $d < \left(1 + \frac{\sigma_1}{\sigma_2}\right) \sqrt{CK}$.

We now consider poGAMBITTS with a nonlinear reward model. In Theorem 2, the complexity measures $\dim_E(\mathcal{F}, T^{-2})$ and $\mathcal{N}(\mathcal{F}, T^{-2}, \|\cdot\|_\infty)$ denote the eluder dimension and the covering number of \mathcal{F} , respectively. Appendix E provides formal definitions of these measures.

Theorem 2. *Let f_1^{off} satisfy (2). For any function class \mathcal{F} with $\hat{m}_2 \in \mathcal{F}$, the Bayesian regret of poGAMBITTS (Algorithm 2) satisfies*

$$BR_T^{\text{po:}\mathcal{F}} \leq \tilde{\mathcal{O}}(\sigma_2 \sqrt{\dim_E(\mathcal{F}, T^{-2}) \log \mathcal{N}(\mathcal{F}, T^{-2}, \|\cdot\|_\infty)} T + T \sqrt{\varepsilon T}).$$

5.2 Regret Analysis for foGAMBITTS

We now present regret guarantees for foGAMBITTS (Algorithm 1) in the linear reward setting, which applies when the agent does not have simulator access to the distributions $\{\xi_{x,a}\}_{(x,a) \in \mathcal{X} \times \mathcal{A}}$.

Theorem 3. *If $m_2(z, x_t; \theta)$ is linear in $z \in \mathbb{R}^d$, the Bayesian regret of foGAMBITTS (Algorithm 1) satisfies*

$$BR_T^{\text{fo:lin}} \leq \tilde{\mathcal{O}}(\sigma_1 \sqrt{CKT} + \sigma_2 d \sqrt{T}).$$

The bound nearly matches the minimax lower bound of $\Omega(\max\{\sigma_1\sqrt{CKT}, \sigma_2\sqrt{dT}\})$, presented in Appendix E. Comparing Theorem 3 and Equation 1 shows foGAMBITTS achieves a sharper regret bound than standard Thompson sampling whenever $d \ll \sqrt{CK}$. When the working treatment representation is low-dimensional and well-specified, the benefits from generalizing across arms outweigh the cost of estimating the treatment distribution. Moreover, the advantage grows when $\sigma_2 \gg \sigma_1$, since generalizing across arms helps more when reward noise dominates. In addition, when $\sigma_1 < 1/\sqrt{CK}$, the regret of foGAMBITTS becomes comparable to that of poGAMBITTS. In such cases, where the generator is sufficiently concentrated, simulator access offers little improvement.

Having established regret guarantees for the GAMBITTS algorithms, we now turn to their empirical evaluation in Section 6.

6 Simulation Results

To evaluate GAMBITTS-based algorithms, we designed a simulation study modeled on the 2023 Intern Health Study (IHS) which was aimed at supporting mental health among medical interns [40]. At each decision point, the agent (*i*) observes user context (location, recent step count), (*ii*) selects a prompt from a finite list to submit to an LLM (Llama 3.1 8.0B), and (*iii*) delivers the generated text response as the intervention [15]. The environment then generates a reward based on select semantic dimensions of the generated text (optimism, severity, formality, clarity, and encouragement). The semantic dimensions were constructed by training one-dimensional variational autoencoders on text generated to vary along each axis, yielding mappings from text to scores used in the reward model (see Appendix G.3 for more details).

While no deployed JITAIs currently integrate LLMs for real-time message generation (and thus no real-world dataset exists for this setting), we calibrate our simulation using empirical distributions from the 2023 IHS to model realistic conditional reward structures. Details on the compute resources, generative model, prompt design, and reward specification are provided in Appendix G, with code and further documentation available on GitHub. Additionally, Appendix F includes further simulations exploring different first- and second-stage variance decompositions, varying the \mathcal{Z} dimensionality d , altering the level of simulation access available to poGAMBITTS, incorporating direct covariate influence on reward, and evaluating performance under nonlinear data-generating mechanisms. All figures are based on 250 Monte Carlo runs per agent, with 95% confidence intervals shown.

6.1 Illustrative Example

We begin with a simple example where the outcome at time t is given by $Y_t = \beta Z_t^{\text{optimism}} + \varepsilon_t$, with Z_t^{optimism} denoting the optimism score of the generated message and $\varepsilon_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma)$. As discussed above, parameters β and σ are derived from IHS data. This setup isolates a single semantic dimension, where we expect GAMBITTS to outperform standard Thompson sampling by leveraging shared reward structure. We compare GAMBITTS algorithms to both contextual and non-contextual standard Thompson sampling agents (StdTS:Contextual and StdTS, respectively); as shown in Figure 2, StdTS performs better here and is used as the baseline in subsequent simulations. As expected, GAMBITTS methods achieve lower regret than standard Thompson sampling, with poGAMBITTS performing especially well in this setting.

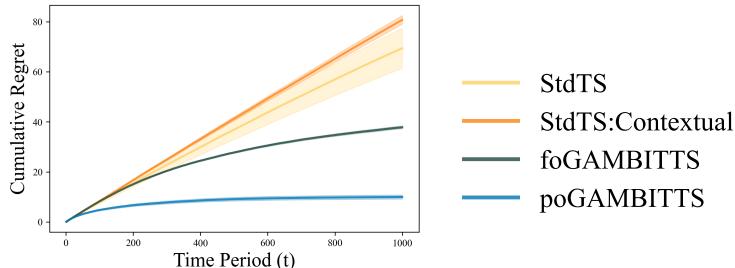


Figure 2: Cumulative Regret Under Single-Dimension Reward Model

6.2 Embedding Misspecification

The assumption that the agent has correctly identified how G influences Y (i.e., that $Z = Z^*$) is quite strong. To assess the impact of misspecifying this mechanism, we revisit the illustrative setting from Section 6.1, but evaluate performance when the agent uses a linear model with an alternative embedding. As discussed in Section 4, we expect GAMBITTS performance to degrade as the working embedding departs from the true reward-relevant dimension (optimism).

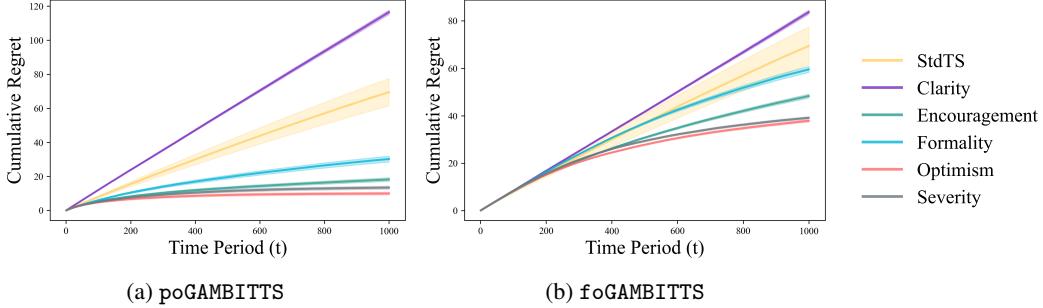


Figure 3: Cumulative Regret Under Treatment Embedding Misspecification

Figure 3 confirms this pattern: GAMBITTS performs well when the working embedding is correlated with the true one, but deteriorates as the correlation declines (see Table G.3.4 in Appendix G). In low-correlation settings, it can even underperform standard Thompson sampling, as policy learning fails when working treatment representations are misaligned with rewards.

6.3 Scaling Number of Arms in a More Complex Reward Structure

We next vary the number of treatment arms to test Corollary 1, which predicts poGAMBITTS regret does not scale with K in linear settings. Furthermore, in this section and throughout Appendix F, we turn to a more complex linear reward structure based on Z^{optimism} , $Z^{\text{formality}}$, and Z^{severity} . As we move to more complex data-generating environments, we introduce an ens-poGAMBITTS agent with an MLP reward model.

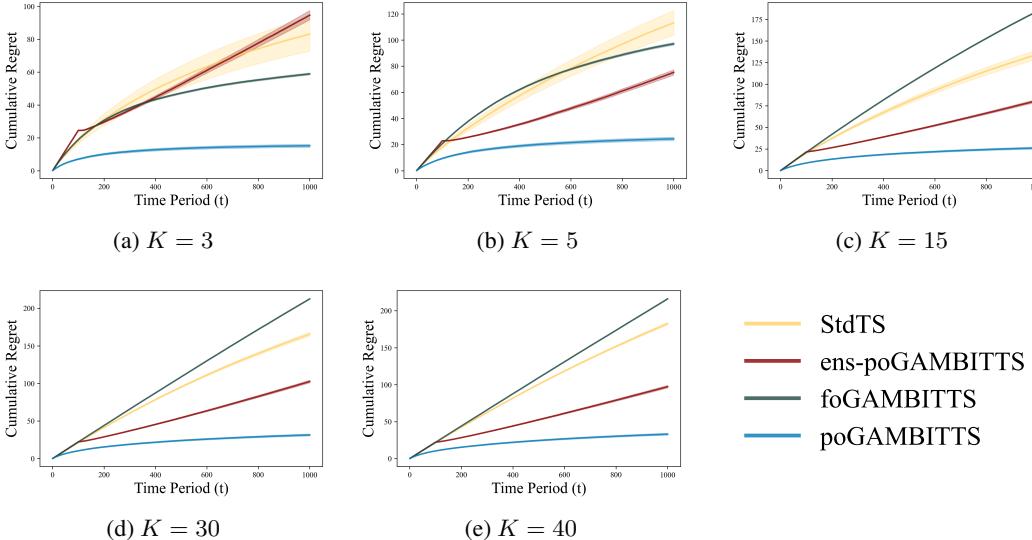


Figure 4: Cumulative Regret for Varying Across Sizes of Action Space

The simulations varying the size of the action set largely aligned with our expectations: the regret of the linear poGAMBITTS algorithm was stable across values of K , and the regret of ens-poGAMBITTS increased slightly with K , but not substantially. This pattern is encouraging, as it suggests that

additional intervention options can be incorporated by increasing K without incurring a large regret penalty, though computational cost scales linearly with K . `foGAMBITTS` performed noticeably worse in this setting, which is expected given that it must estimate the distribution of Z for each (x, a) pair. A more natural comparator for `foGAMBITTS` is the contextual bandit, and we examine this comparison in Appendix F.

7 Discussion and Future Work

This paper introduces GAMBITTS for online learning in generator-mediated bandit environments, showing promise for a synergistic relationship between classical bandit methods and modern advances in generative modeling. By formalizing settings in which actions (e.g., prompts) produce stochastic treatments (e.g., generated responses), this framework opens new opportunities for personalization in domains such as marketing, mHealth, and education, where treatments cannot be easily enumerated and real-time generation can yield more deeply-tailored interventions.

This work opens up a wide range of directions for future research. One immediate area concerns the structure and specification of the working treatment embedding. As shown in Section 6, GAMBITTS can still perform well under embedding misspecification, and flexible reward models can help mitigate this issue. Still, a natural extension would be to learn the embedding online, (e.g., through online sufficient dimensionality reduction). Another direction involves moving beyond a static generator. While this manuscript assumes a fixed model, fine-tuning the generator based on observed outcomes could offer new opportunities for intervention design. Finally, this manuscript assumes that the agent always delivers the generated response to the environment. While this may be appropriate in some applications, in others it may be essential to ensure that any content sent to users is safe. In such cases, intervention designers may wish to incorporate safety checks or content filters prior to delivery. Incorporating these constraints into the generator-mediated bandit framework may require new GAMBITTS variants and accompanying theoretical analysis.

Acknowledgments

The authors thank researchers at d3center for their comments in the early stages of this work; in particular: Daniel Almirall, Walter Dempsey, Mason Ferlic, Easton Huch, Wenchu Pan, Yao Song, and Shiyu Zhang. We also thank Jordy Berne, Liza Levina, and Abhiti Mishra for their comments on later drafts of the manuscript, as well as Elena Frank and Srijan Sen, director and principal investigator of the Intern Health Study, for their support and coordination. Lastly, Gabriel Durham and Ambuj Tewari acknowledge support from the NIH via grants 5R01DA039901 and P50DA054039, and 5R01MH131617 and 5U01MH136025, respectively.

References

- [1] Jason Ansel, Edward Z. Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, Peter Bell, David Berard, Evgeni Burovski, Geeta Chauhan, Anjali Chourdia, et al. Pytorch 2: Faster machine learning through dynamic Python bytecode transformation and graph compilation. In *ASPLOS* (2), pages 929–947, 2024.
- [2] Mirza Mansoor Baig, Chris Hobson, Hamid GholamHosseini, Ehsan Ullah, and Shereen Afifi. Generative AI in improving personalized patient care plans: Opportunities and barriers towards its wider adoption. *Applied Sciences*, 14(23):10899, November 2024.
- [3] Hamsa Bastani and Mohsen Bayati. Online decision making with high-dimensional covariates. *Operations Research*, 68(1):276–294, January 2020.
- [4] Audrey Boruvka, Daniel Almirall, Katie Witkiewitz, and Susan A. Murphy. Assessing time-varying causal effect moderation in mobile health. *Journal of the American Statistical Association*, 113(523):1112–1121, July 2018.
- [5] Djallel Bouneffouf and Raphael Feraud. Survey: Multi-armed bandits meet large language models, 2025.
- [6] Clément Chadebec, Louis J. Vincent, and Stéphanie Allassonnière. Pytha: Unifying generative autoencoders in Python – a benchmarking use case, 2023.
- [7] Lichang Chen, Juhai Chen, Tom Goldstein, Heng Huang, and Tianyi Zhou. InstructZero: Efficient instruction optimization for black-box large language models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235, pages 6503–6518. PMLR, 21–27 Jul 2024.
- [8] Zekai Chen, Po-Yu Chen, and Francois Buet-Golfouse. Online personalizing white-box LLMs generation with neural bandits. In *Proceedings of the 5th ACM International Conference on AI in Finance*, ICAIF ’24, page 711–718. ACM, November 2024.
- [9] Riccardo Della Vecchia and Debabrota Basu. Stochastic online instrumental variable regression: Regrets for endogeneity and bandit feedback. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(15):16190–16198, April 2025.
- [10] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding, 2018.
- [11] Naoki Egami, Christian J. Fong, Justin Grimmer, Margaret E. Roberts, and Brandon M. Stewart. How to make causal inferences using texts. *Science Advances*, 8(42), October 2022.
- [12] Khaled Eldowa, Nicolò Cesa-Bianchi, Alberto Maria Metelli, and Marcello Restelli. Information capacity regret bounds for bandits with mediator feedback. *J. Mach. Learn. Res.*, 25(1), January 2024.
- [13] Lars Fritzsche, Elena Frank, Laura Thomas, Amy Bohnert, and Srijan Sen. Implementing large language models for tailored mental health messaging in the COMPASS study, 2024.
- [14] Susobhan Ghosh, Yongyi Guo, Pei-Yao Hung, Lara Coughlin, Erin Bonar, Inbal Nahum-Shani, Maureen Walton, and Susan Murphy. ReBandit: Random effects based online RL algorithm for reducing cannabis use. *IJCAI (U. S.)*, 2024:7278–7286, August 2024.

- [15] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, et al. The Llama 3 herd of models, 2024.
- [16] David Haag, Devender Kumar, Sebastian Gruber, Mahdi Sareban, Gunnar Treff, Josef Niebauer, Christopher Bull, and Jan David Smeddinck. The last JITAI? The unreasonable effectiveness of large language models in issuing just-in-time adaptive interventions: Fostering physical activity in a prospective cardiac rehabilitation setting, 2024.
- [17] Jawad Hazime and Junaid Farooq. Evaluation of LLM powered agentic AI for solving multi-arm bandit problems. In *2025 IEEE International Conference on Omni-layer Intelligent Systems (COINS)*, page 1–6. IEEE, August 2025.
- [18] Ji He, Jianshu Chen, Xiaodong He, Jianfeng Gao, Lihong Li, Li Deng, and Mari Ostendorf. Deep reinforcement learning with a natural language action space. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, page 1621–1630. Association for Computational Linguistics, 2016.
- [19] Xinyu Hu, Min Qian, Bin Cheng, and Ying Kuen Cheung. Personalized policy learning using longitudinal mobile health data. *Journal of the American Statistical Association*, 116(533):410–420, August 2020.
- [20] Easton Knight Huch, Jieru Shi, Madeline R Abbott, Jessica R Golbus, Alexander Moreno, and Walter H. Dempsey. RoME: A robust mixed-effects bandit algorithm for optimizing mobile health interventions. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, 2024.
- [21] Kosuke Imai and Kentaro Nakamura. Causal representation learning with generative artificial intelligence: Application to texts as treatments, 2024.
- [22] Theodore J. Iwashyna and Edward H. Kennedy. Instrumental variable analyses. Exploiting natural randomness to understand causal mechanisms. *Annals of the American Thoracic Society*, 10(3):255–260, June 2013.
- [23] Lorenzo J. James, Laura Genga, Barbara Montagne, Muriel A. Hagenaars, and Pieter M. E. Van Gorp. Toward scalable content generation for gamified mHealth interventions: The evaluation of LLM-generated goals on user engagement. In *Entertainment Computing – ICEC 2024*, pages 94–107, Cham, 2025. Springer Nature Switzerland.
- [24] Jussi S. Jauhainen and Agustín Garagorry Guerra. Generative AI and education: Dynamic personalization of pupils’ school learning material with ChatGPT. *Frontiers in Education*, 9, November 2024.
- [25] Nathan Kallus. Instrument-armed bandits. In *Proceedings of Algorithmic Learning Theory*, volume 83, pages 529–546. PMLR, April 2018.
- [26] Jong-Min Kim. Llm-guided ensemble learning for contextual bandits with copula and gaussian process models. *Mathematics*, 13(15):2523, August 2025.
- [27] Akshay Krishnamurthy, Keegan Harris, Dylan J Foster, Cyril Zhang, and Aleksandrs Slivkins. Can large language models explore in-context? In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, 2024.
- [28] Nir Kshetri, Yogesh K. Dwivedi, Thomas H. Davenport, and Niki Panteli. Generative artificial intelligence in marketing: Applications, opportunities, challenges, and research agenda. *International Journal of Information Management*, 75, April 2024.
- [29] Branislav Kveton, Yi Liu, Johan Matteo Kruijssen, and Yisu Nie. Non-compliant bandits. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*, page 1138–1147, 2023.
- [30] Finnian Lattimore, Tor Lattimore, and Mark D. Reid. Causal bandits: Learning good interventions via causal inference. In *Advances in Neural Information Processing Systems 29 (NIPS 2016)*, page 1189–1197, 2016.

- [31] Tor Lattimore and Csaba Szepesvari. *Bandit algorithms*. Cambridge University Press, 2020.
- [32] Yingkai Li, Yining Wang, and Yuan Zhou. Nearly minimax-optimal regret for linearly parameterized bandits. In *Conference on Learning Theory*, pages 2173–2174. PMLR, 2019.
- [33] Peng Liao, Kristjan Greenewald, Predrag Klasnja, and Susan Murphy. Personalized HeartSteps: A reinforcement learning algorithm for optimizing physical activity. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 4(1):1–22, March 2020.
- [34] Xiaoqiang Lin, Zhaoxuan Wu, Zhongxiang Dai, Wenyang Hu, Yao Shu, See-Kiong Ng, Patrick Jaillet, and Bryan Kian Hsiang Low. Use your INSTINCT: INSTruction optimization for LLMs usIng Neural bandits Coupled with Transformers. In *Proceedings of the 41st International Conference on Machine Learning*, ICML’24. JMLR.org, 2024.
- [35] Ziyi Liu, Idan Attias, and Daniel M. Roy. Causal bandits: The Pareto optimal frontier of adaptivity, a reduction to linear bandits, and limitations around unknown marginals. In *Proceedings of the 41st International Conference on Machine Learning*. PMLR, 2024.
- [36] Xiuyuan Lu and Benjamin Van Roy. Ensemble sampling. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, page 3260–3268, 2017.
- [37] Yangyi Lu, Amirhossein Meisami, Ambuj Tewari, and Zhenyu Yan. Regret analysis of bandit problems with causal background knowledge. In *Conference on Uncertainty in Artificial Intelligence*, 2019.
- [38] Aурghya Maiti, Vineet Nair, and Gaurav Sinha. A causal bandit approach to learning good atomic interventions in presence of unobserved confounders. In *Proceedings of the Thirty-Eighth Conference on Uncertainty in Artificial Intelligence*, volume 180, pages 1328–1338. PMLR, August 2022.
- [39] Inbal Nahum-Shani, Shawna N Smith, Bonnie J Spring, Linda M Collins, Katie Witkiewitz, Ambuj Tewari, and Susan A Murphy. Just-in-time adaptive interventions (JITAIs) in mobile health: Key components and design principles for ongoing health behavior support. *Annals of Behavioral Medicine*, 52(6):446–462, December 2017.
- [40] Timothy NeCamp, Srijan Sen, Elena Frank, Maureen A Walton, Edward L Ionides, Yu Fang, Ambuj Tewari, and Zhenke Wu. Assessing real-time moderation for developing adaptive mobile health interventions for medical interns: Micro-randomized trial. *Journal of Medical Internet Research*, 22(3):e15033, March 2020.
- [41] Allen Nie, Yi Su, Bo Chang, Jonathan Lee, Ed Chi, Quoc Le, and Minmin Chen. EVOLvE: Evaluating and optimizing LLMs for in-context exploration. In *Proceedings of the 42nd International Conference on Machine Learning*, ICML’25, 2025.
- [42] Ollama. Ollama (version 0.3.3). <https://ollama.com>, 2024.
- [43] Reid Pryzant, Dallas Card, Dan Jurafsky, Victor Veitch, and Dhanya Sridhar. Causal effects of linguistic properties. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, 2021.
- [44] Daniel Russo and Benjamin Van Roy. Eluder dimension and the sample complexity of optimistic exploration. *Advances in Neural Information Processing Systems 26 (NIPS 2013)*, 26, 2013.
- [45] Rajat Sen, Karthikeyan Shanmugam, Murat Kocaoglu, Alex Dimakis, and Sanjay Shakkottai. Contextual Bandits with Latent Confounders: An NMF Approach. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54, pages 518–527. PMLR, 20–22 Apr 2017.
- [46] Chengshuai Shi, Kun Yang, Zihan Chen, Jundong Li, Jing Yang, and Cong Shen. Efficient prompt optimization through the lens of best arm identification. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, 2024.
- [47] Aleksandrs Slivkins. Introduction to multi-armed bandits, 2019.

- [48] Oleg Stavitskii, Kyrylo Bulatsev, Philipp Petrenko, Dmitry Bezugly, Evgeny Gurzhiy, and Dmitry Evgrafov. System and method for creating a personalized user environment. U.S. Patent 11,275,350 B2, 2022. <https://patents.justia.com/patent/11275350>.
- [49] Andrew Stirn and Tony Jebara. Thompson sampling for noncompliant bandits, 2018.
- [50] Brian Sustoletto. Deceptively simple yet profoundly impactful: Text messaging interventions to support health. *Journal of Medical Internet Research*, 26, August 2024.
- [51] Ambuj Tewari and Susan A. Murphy. *From Ads to Interventions: Contextual Bandits in Mobile Health*, page 495–517. Springer International Publishing, 2017.
- [52] Sabina Tomkins, Peng Liao, Predrag Klasnja, and Susan Murphy. IntelligentPooling: Practical Thompson sampling for mhealth. *Machine Learning*, 110(9):2685–2727, June 2021.
- [53] Victor Veitch, Dhanya Sridhar, and David M. Blei. Adapting text embeddings for causal inference. In *Conference on Uncertainty in Artificial Intelligence*, 2019.
- [54] Jiangxing Wang, Jiachen Li, Xiao Han, Deheng Ye, and Zongqing Lu. Language model adaption for reinforcement learning with natural language action space. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, page 1620–1634. Association for Computational Linguistics, 2024.
- [55] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierrick Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, and Jamie Brew. HuggingFace’s transformers: State-of-the-art natural language processing. *CoRR*, 2019.
- [56] Nuoya Xiong and Wei Chen. Combinatorial pure exploration of causal bandits. In *The 11th International Conference on Learning Representations*, 2023.
- [57] Akihiro Yabe, Daisuke Hatano, Hanna Sumita, Shinji Ito, Naonori Kakimura, Takuro Fukunaga, and Ken-ichi Kawarabayashi. Causal bandits with propagating inference. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80, pages 5512–5520. PMLR, 10–15 Jul 2018.
- [58] Jingwen Zhang, Yifang Chen, and Amandeep Singh. Causal bandits: Online decision-making in endogenous settings. In *NeurIPS 2022 Workshop on Causality for Real-world Impact*, 2022.
- [59] Lutong Zou, Ziping Xu, Daiqi Gao, and Susan Murphy. Causal directed acyclic graph-informed reward design, 2025.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction describe the proposed generator-mediated bandit framework and the Thompson sampling-based algorithm (GAMBITTS) (these are presented in Sections 3 and 4, respectively). They also summarize the theoretical regret bounds and empirical performance improvements over standard baselines, aligning with the theoretical results in Section 5 and the empirical results in Section 6.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss limitations related to the reliance on a correctly specified working embedding and examine robustness to misspecification in Section 6.2. Additionally, as noted in Section 5, the performance of poGAMBITTS depends on the quality of the first-stage treatment approximation, and the theoretical results are derived under simplified conditions, including a fixed context space and treatment model.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best

judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: We explain the assumptions for our main body results in Section 5. We include proofs in Appendix E.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: We include algorithmic details in Section 4, simulation specifications in Appendix G, and we include the code necessary to reproduce our results in the supplementary material.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.

- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We provide code in the supplementary materials with instructions for reproducing the simulation results reported in Section 6 and Appendix F. This code will be made publicly available on GitHub upon publication; the link is omitted here to preserve anonymity.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We describe the simulation setup and implementation details in Appendix G, including data-generating mechanisms and hyperparameters. We also include code in the supplementary materials to ensure results can be reproduced exactly.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: As discussed in Section 6, we report 95% confidence intervals for Monte Carlo estimates of cumulative reward. These intervals are computed across 250 independent runs per agent and reflect uncertainty due to randomization in the environment and learning process.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We present this information in Appendix G.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research complies with the NeurIPS Code of Ethics. We have, to the best of our ability, followed best practices in licensing, reproducibility, and responsible use of external models and datasets, and we address potential limitations and risks in Section 7.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Yes. We discuss broader social impacts in Appendix H, including potential benefits in areas like education and mHealth. We also note risks associated with deploying LLM-generated text in sensitive settings and highlight future directions for adapting the framework to mitigate those risks (see Section 7).

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: We do not release any new pretrained models or datasets. However, we discuss potential risks and implications of deploying algorithms like ours in sequential decision-making settings in Section 7.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We use Llama 3.1 to generate text-based treatments in our simulations and cite the official technical report. We used the model in accordance with its non-commercial research license, and all other external assets are properly credited and used within their stated terms.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The paper introduces a simulation framework and a pool of LLM-generated text-based treatments. These are included in the supplementary materials along with code and documentation describing how they were generated and used in the experiments.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: We did not conduct any experiments involving human subjects or crowdsourcing. However, some simulation parameters were based on prior published studies involving human participants, which we cite appropriately.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: Some simulation parameters were derived from prior published studies with human subjects, which we cite. This paper does not involve new human subjects research and did not require IRB approval.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs are used to generate a pool of text-based treatments, which serve as the basis for our simulation experiments. While the LLM is not part of the learning algorithm itself, it plays a role in motivating the problem setup (described in Section 3) and its use in the simulation experiments is described in Section 6 and Appendix G.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

A Further Related Work

Eldowa et al. [2024] study *bandits with mediator feedback* from an information-theoretic perspective, broadly considering settings where the effect of an action is funneled through a stochastic mediator [12].² Among works on bandits with mediator feedback, our approach is most closely related to those that draw connections to instrumental variables and noncompliance. Zhang et al. [2022] and Della Vecchia and Basu [2025] adopt an econometric perspective, developing OFUL-based algorithms for bandit problems with endogenous covariates. While both consider continuous treatments, they do not focus on high-dimensional treatment spaces or settings in which treatments can be generated offline without interacting with the environment [58, 9]. Moreover, these works emphasize interpretable parameter estimation, particularly in the structural model of the reward, whereas our primary focus is on predictive modeling for policy learning. This distinction is important in generative settings, where treatments (e.g., LLM-generated text) are difficult to interpret directly, and causal inference is used to support adaptive decision-making rather than uncover underlying mechanisms.

Other related efforts include Stirn and Jebara [2018], which introduces Thompson sampling in the context of noncompliance, assuming a shared, discrete action and treatment space [49]. Kveton et al. [2023] also study noncompliance and employ a Thompson sampling-based approach, aiming to identify actions with the highest compliance-weighted mean reward, rather than modeling the full stochastic treatment-generation process [29].

The generator-mediated bandit setup most closely resembles the noncompliance-inspired *instrument-armed bandit* (IAB) framework of Kallus [2018], where each arm pull represents the choice of an instrument, and the mediated bandit environment of Zou et al. [2025] [25, 59]. One key departure in our setting is the presence of a stochastic generator that produces the treatment, which leads us to focus on questions of treatment representation, generator access, and nonlinear outcome modeling, areas not emphasized in the IAB or other mediated bandit frameworks.

In contrast to Kallus’ focus, our setting does not require the treatment space \mathcal{G} to match the action space \mathcal{A} . While assuming a shared action and treatment space is natural for modeling noncompliance, where the intended and realized actions are ideally aligned, our setting explicitly separates the two. Here, the distinction is fundamental rather than incidental, reflecting the design goal of producing personalized responses rather than enforcing direct action execution.

Like Zou et al. [2025], we consider a general framework for learning in mediated bandit environments. However, motivated by the complex nature of generative outputs, we focus on flexible, potentially nonlinear models for both the mediator and the reward. In contrast, Zou et al. frame their approach more explicitly as a surrogate-reward method,³ emphasizing linear models and aiming for robustness in cases where the action may influence the reward directly, outside the mediation pathway.⁴

The works discussed above present algorithms that can be broadly understood as instances of a general *noisy action* framework, where the environment observes a stochastic transformation of the agent’s chosen action.⁵ From this perspective, GAMBITTS represents a specific instantiation of a noisy-action Thompson sampler. While GAMBITTS is motivated by the GAMBIT framework, with high-dimensional and continuous treatment spaces, taking $h = I$ (as defined in Section 3) recovers algorithms suited to the econometric and psychometric mediated bandit environments discussed above. Further taking $\mathcal{G} = \mathcal{A}$ yields algorithms appropriate for noncompliance settings.

In addition to bandit methods with comparable structure, described above, another line of research has explored intersections between LLMs and bandits. Although these approaches address different statistical questions, they nonetheless relate to our setting through their shared focus on combining LLMs with bandit methods. Within this broader category, the most directly related line of work concerns using bandits for prompt selection, where prompts are treated as arms and evaluated based on task-specific performance. In contrast with our approach, much of this literature considers settings with deterministic rewards once the LLM output is observed, or with white-box access to the model

²Their analysis focuses on discrete mediators and introduces an EXP4-based algorithm.

³We note that GAMBITTS can also be viewed in this way, as discussed in Section 5.

⁴In our motivating application, the action affects the outcome only through the generated treatment, making full mediation a reasonable assumption.

⁵We use the term *noisy action* to emphasize that, from the environment’s perspective, the observed treatment is a noisy representation of the agent’s intended action.

[7, 8, 34, 46]. Another line of recent research examines LLMs themselves as the decision-making agents in multi-armed bandit [27, 17, 26, 41]. For a broader overview of work at the intersection of LLMs and bandits, see Bouneffouf and Feraud [2025] [5].

In the mHealth literature, researchers have recently begun exploring how LLMs can support intervention design, motivated by the the parallel rise of mobile interventions interventions and generative models. Haag et al. [2024] compared the quality of JITAI intervention text generated by GPT-4 to that authored by both domain experts and non-experts. Their findings showed that LLM-generated content performed favorably compared with human-authored text [16]. Additionally, James et al. [2024] conducted a randomized trial comparing LLM- and human-generated goals in a gamified mHealth application, evaluating their impact on participant engagement. The study found no significant difference in engagement between the two groups, suggesting that LLMs may offer a viable and scalable alternative for content generation in such settings [23]. However, as the goals were pre-written and not personalized in real time; the study did not evaluate the use of LLMs for on-the-fly message generation or adaptive personalization. Lastly, the IHS-COMPASS studies used LLMs to generate JITAI intervention text in recent deployments. However, as in James et al. [2024], these messages were pre-written rather than generated on-the-fly [13].

B Practical Variations

B.1 Restricting Context for the Generator

As discussed in Section 3, in many practical applications, the agent may restrict which features of X_t are passed to the generator, selecting a subset $X_t^G \subseteq X_t$ to include in the query. While some components of X_t may be useful for outcome modeling, they may not be appropriate for inclusion in the prompt sent to the generator. For example, a variable like weight could improve reward prediction but would likely be excluded from the generator input for ethical or design reasons. The corresponding causal structure is shown in Figure 5.

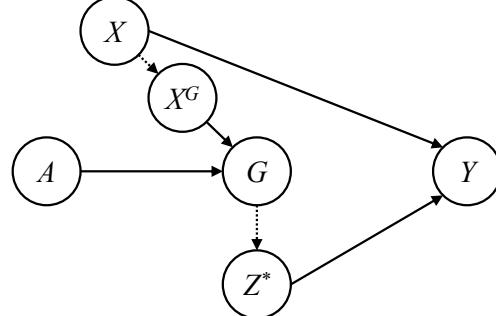


Figure 5: Generator-Mediated Bandit Causal Structure (Restricted Context in Query)

- Notes:*
- (i) Dotted lines represent deterministic relationships.
 - (ii) Arrows from X to A are omitted to represent the data-generating process rather than the decision logic.
 - (iii) A : Action/Query; X : Full Context; X^G : Context Sent to Generator; G : Generated Response; Z^* : True Response Embedding; Y : Reward.

The methods and analyses in this manuscript change only superficially when $X_t^G \neq X_t$. In this case, the treatment model introduced in Section 4, $f_1(z; A_t, X_t, \theta_1)$, is replaced by $f_1(z; A_t, X_t^G, \theta_1)$. Additionally, all algorithms that invoke f_1 inherit this modification.

B.2 Optional Prompting Extension

As discussed in Section 4, mHealth interventions frequently begin with a decision of whether to deliver any message at all. GAMBITTS can be naturally extended to incorporate this structure. In this variant, a primary agent first decides whether to act (e.g., to send a message). A GAMBITTS-based agent then operates only at the time points where the primary agent chooses to act. Algorithm 3 outlines this extension.

Algorithm 3 GAMBITTS with Optional Prompting

Inputs: Full history $\mathcal{D}_t^{(1)}$, GAMBITTS history $\mathcal{D}_t^{(2)}$, prompting policy $\mu^{(1)}$, GAMBITTS policy $\mu^{(2)}$, context x_t .

- 1: **Sample** $R_t \sim \mu^{(1)}(\mathcal{D}_t^{(1)}, x_t)$ according to Agent 1.
- 2: **If** $R_t = 1$ **then**
- 3: Select query $A_t \sim \mu^{(2)}(\mathcal{D}_t^{(2)}, x_t)$ according to Agent 2.
- 4: Update $\mathcal{D}_{t+1}^{(2)} \leftarrow \mathcal{D}_t^{(2)} \cup \{(x_t, A_t, Z_t, Y_t)\}$.
- 5: Update policy $\mu^{(2)}$.
- 6: **Update** $\mathcal{D}_{t+1}^{(1)} \leftarrow \mathcal{D}_t^{(1)} \cup \{(x_t, R_t, Y_t)\}$.
- 7: **Update** policy $\mu^{(1)}$.

C Embedding Misspecification

As discussed in Section 4, GAMBITTS performance hinges on the construction of working embedding $Z = h(G)$, with the agent modeling $Y | Z, X$ in place of $Y | G, X$. Section 6.2 supports this notion, showing GAMBITTS is robust to mild misspecification but suffers under severe misspecification. With this phenomenon in mind, this appendix discusses two approaches aimed at improving robustness to embedding misspecification.

One option for promoting robustness is to construct a hedged reward model. In particular, given any initial embedding construct Z and model m_2 , an analyst can consider the modified form:

$$\tilde{m}_2(A_t, Z_t, X_t; \beta^*, \theta_2) = \sum_{a \in \mathcal{A}} \beta^a \mathbb{1}_{A_t=a} + m_2(Z_t, X_t; \theta_2).$$

When the goal is online learning of the optimal action A , this structure offers a hedge against misspecification of Z and m_2 . In particular, if m_2 fails to capture signal, its estimated contribution to \tilde{m}_2 will shrink towards zero. In that case, the \tilde{m}_2 model asymptotically behaves like a standard Thompson sampling agent that learns a reward distribution over actions directly. Conversely, if Z and m_2 are well-specified, then the action A_t provides no additional information beyond what is already captured through $m_2(Z_t, X_t; \theta_2)$. In this case, the estimated β^* terms will shrink to zero and \tilde{m}_2 will converge to m_2 . However, the agent must first learn that the embedding carries signal, potentially delaying efficient generalization early in learning.

Figure 6 revisits the analysis presented in Figure 3 of Section 6.2, now including hedged models. All hedged poGAMBITTS agents perform similarly, as do all hedged foGAMBITTS agents. While every hedged GAMBITTS variant outperforms standard Thompson sampling in this setting, the well-specified models no longer exhibit the performance advantage they had in the original (unhedged) analysis. This supports the hypothesis that hedging reduces downside risk but limits gains under well-specified embeddings.

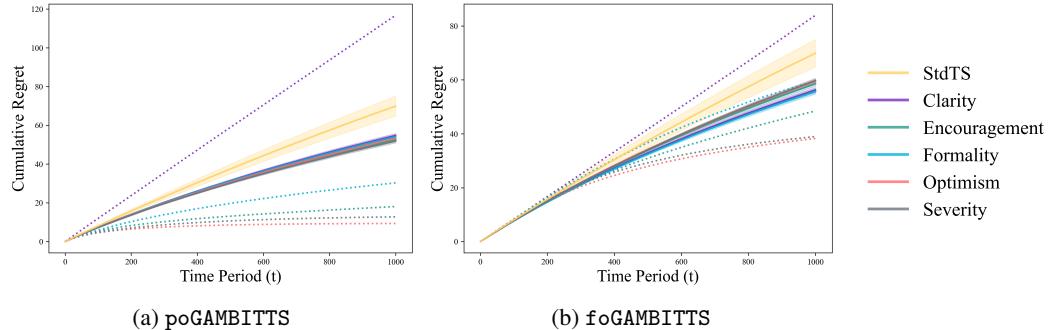


Figure 6: Cumulative Regret Under Embedding Misspecification - Original vs Hedged Approaches
Note: Dotted lines represent cumulative regret under original (unhedged) approaches.

A second option is to maintain a set of candidate reward models, each potentially using different treatment embeddings. Periodically, the agent can evaluate these models based on predictive performance and switch to the one that best predicts observed outcomes. This allows the algorithm to recover from poor initial choices of Z (and m_2 more broadly) without requiring strong prior knowledge about which embedding will perform best.

Improving robustness to embedding misspecification also points to several directions for future work. For example, extending the candidate model approach by developing a corraling bandit, where a master algorithm selects among reward models, could allow the agent to adaptively shift across models and tailor the choice of embedding to the individual. Similarly, as discussed in Section 7, learning individualized treatment embeddings through methods like online sufficient dimensionality reduction would mitigate misspecification risk.

D Ensemble-Based Algorithms

As in Lu and Van Roy [2017], given a second-stage reward model parametrized by θ_2 , ensemble-based GAMBITTS approaches require initially drawing M_{ens} sets of parameters $\tilde{\theta}_{2,1}^1, \dots, \tilde{\theta}_{2,1}^{M_{ens}}$ from a pre-determined initialization distribution. Algorithms 4 and 5 then demonstrate how to select action a_t for $t = 1, \dots, T$. As in Algorithms 1 and 2, an analyst may approximate the integrals below via Monte Carlo.

Steps 5 and 4 in Algorithms 4 and 5 (respectively) require updating model parameters as described in Lu and Van Roy [2017]. These updates involve perturbations of observed reward to promote exploration [36].

Algorithm 4 Ensemble-Based Fully Online GAMBITTS (ens-foGAMBITTS)

Inputs: Data \mathcal{D}_t , prior π_1 , models f_1 , m_2 , current context x_t , reward model parameters $\left\{ \tilde{\theta}_{2,t}^m \right\}_{m=1}^{M_{ens}}$.

- 1: **Sample** $\theta_{1,t} \sim P_1(\theta_1 | \mathcal{D}_t)$
- 2: **Sample** $j \sim \mathcal{U}(\{1, \dots, M_{ens}\})$.
- 3: **For** each $a \in \mathcal{A}$, calculate

$$\mathbb{E}_{\theta_{1,t}, \tilde{\theta}_{2,t}^j} [Y | a, x_t] = \int_{\mathcal{Z}} m_2(z, x_t; \tilde{\theta}_{2,t}^j) f_1(z; a, x_t, \theta_{1,t}) dz.$$

- 4: **Set** $A_t = \arg \max_{a \in \mathcal{A}} \mathbb{E}_{\theta_{1,t}, \tilde{\theta}_{2,t}^j} [Y | a, x_t]$.
 - 5: **Given** (z_t, x_t, y_t) , update $\tilde{\theta}_{2,t+1}^1, \dots, \tilde{\theta}_{2,t+1}^{M_{ens}}$.
-

Algorithm 5 Ensemble-Based Partially Online GAMBITTS (ens-poGAMBITTS)

Inputs: Data \mathcal{D}_t , model m_2 , current context x_t , pretrained model f_1^{off} , reward model parameters $\left\{ \tilde{\theta}_{2,t}^m \right\}_{m=1}^{M_{ens}}$.

- 1: **Sample** $j \sim \mathcal{U}(\{1, \dots, M_{ens}\})$.
- 2: **For** each $a \in \mathcal{A}$, calculate

$$\mathbb{E}_{\tilde{\theta}_{2,t}^j} [Y | a, x_t] = \int_{\mathcal{Z}} m_2(z, x_t; \tilde{\theta}_{2,t}^j) f_1^{off}(z; a, x_t) dz.$$

- 3: **Set** $A_t := \arg \max_{a \in \mathcal{A}} \mathbb{E}_{\tilde{\theta}_{2,t}^j} [Y | a, x_t]$.
 - 4: **Given** (z_t, x_t, y_t) , update $\tilde{\theta}_{2,t+1}^1, \dots, \tilde{\theta}_{2,t+1}^{M_{ens}}$.
-

E Further Theoretical Discussion

In this appendix, we provide proofs for the results discussed in Section 5, along with additional theoretical developments related to Thompson sampling in the GAMBITTS framework. The appendix is organized as follows:

- Appendix E.1: Proof of the regret upper bound presented in Section 5 for standard Thompson sampling in a GAMBIT environment.
- Appendix E.2: Proofs of the regret upper bounds presented in Section 5 for poGAMBITTS-based algorithms, as well as a bound on the number of samples needed for the offline (empirical) treatment distribution to approximate the true distribution.
- Appendix E.3: Proof of the regret upper bound presented in Section 5 for foGAMBITTS.
- Appendix E.4: Statement and proof of a regret lower bound for foGAMBITTS.

E.1 Upper Regret Bound for Standard Thompson Sampling

In this section, we present a Bayesian regret bound of a standard Thompson sampling algorithm that models the problem as a conventional multi-armed contextual bandit over $\mathcal{X} \times \mathcal{A}$, ignoring the generator-mediation structure. Our analysis adapts the proof from Chapter 36.1 of Lattimore and Szepesvári [2020] [31].

Claim 1. *The Bayesian regret of the standard Thompson sampling agent that models the problem as a multi-armed contextual bandit over $\mathcal{X} \times \mathcal{A}$ is bounded by*

$$BR_T^{\text{standard}} \leq \mathcal{O}((\sigma_1 + \sigma_2)\sqrt{CKT \log T} + BCK)$$

where B is a uniform bound on the mean reward such that $|m_2(z, x; \theta)| \leq B$.

Proof. Recall that the reward received when the agent selects action $a \in \mathcal{A}$ under context $x \in \mathcal{X}$ is given by $m_2(Z, x; \theta_2) + \eta$, where η is σ_2 -subgaussian and $Z \in \mathbb{R}^d$ is drawn from the distribution $\xi_{x,a}^{\theta_1}$.⁶ As discussed in Section 5, we assume that $m_2(Z, x; \theta_2)$ is σ_1 -subgaussian. It follows that the overall reward for a given context-action pair (x, a) is $(\sigma_1 + \sigma_2)$ -subgaussian.

We use $m(x, a)$ to denote the mean reward for the context-action pair (x, a) ; i.e.,

$$m(x, a) := \int_{\mathcal{Z}} m_2(z, x; \theta_2) f_1(z; x, a, \theta_1) dz.$$

Since we assume that the mean reward is bounded by $|m_2(z, x; \theta_2)| \leq B$, it follows that $|m(x, a)| \leq B$.

Now, we adapt the Bayesian regret analysis for a standard Thompson sampling algorithm for multi-armed bandit settings to the *contextual* multi-armed bandit setting where the context is assumed to be drawn i.i.d. every round.

Well, for each $a \in \mathcal{A}$, $x \in \mathcal{X}$ and $t = 1, \dots, T$, define

$$U_t(x, a) := \text{clip}_{[-B, B]} \left(\hat{m}_{x,a}(t-1) + (\sigma_1 + \sigma_2) \sqrt{\frac{2 \log(1/\delta)}{1 \vee N_{x,a}(t-1)}} \right),$$

where $N_{x,a}(t-1) = \sum_{s=1}^{t-1} \mathbb{1}_{\{X_s=x, A_s=a\}}$ is the count of the pair (x, a) up to time $t-1$ and $\hat{m}_{x,a}(t-1)$ denotes the empirical estimate of the reward of action a under context x based on observations up to time $t-1$. If $N_{x,a}(t-1) = 0$, we set $\hat{m}_{x,a}(t-1) = 0$.

Let E be the event that, for all $t = 1, \dots, T$ and $(x, a) \in \mathcal{X} \times \mathcal{A}$, the empirical estimate satisfies

$$|\hat{m}_{t-1}(x, a) - m(x, a)| \leq (\sigma_1 + \sigma_2) \sqrt{\frac{2 \log(1/\delta)}{1 \vee N_{x,a}(t-1)}}.$$

⁶Recall $\xi_{x,a}^{\theta_1}$ has density $f_1(z; x, a, \theta_1)$.

Under this setup, the standard Hoeffding concentration bound gives $\mathbb{P}[E^c] \leq 2CKT\delta$ where $C = |\mathcal{X}|$ and $K = |\mathcal{A}|$.

Let $\mathcal{F}_t = \sigma(X_1, A_1, Y_1, \dots, X_t, A_t, Y_t, X_{t+1})$ be the σ -algebra generated by the interaction sequence by the end of time t (and the following context at time $t+1$). Note that $U_t(x, a)$ is \mathcal{F}_{t-1} -measurable. Using the fact that, under Thompson sampling, the conditional distributions of $A_t^* = \arg \max_a m(X_t, a)$ and A_t given \mathcal{F}_{t-1} are identical, we obtain

$$\begin{aligned} \text{BR}_T &= \mathbb{E} \left[\sum_{t=1}^T (m(X_t, A_t^*) - m(X_t, A_t)) \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \mathbb{E} [m(X_t, A_t^*) - m(X_t, A_t) | \mathcal{F}_{t-1}] \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T (m(X_t, A_t^*) - U_t(X_t, A_t^*)) + \sum_{t=1}^T (U_t(X_t, A_t) - m(X_t, A_t)) \right]. \end{aligned}$$

On the event E^c , the terms inside the expectation are bounded by $2BT$. On the event E , the first sum is negative and the second sum is bounded by

$$\begin{aligned} &\mathbb{1}_{\{E\}} \sum_{t=1}^T (U_t(X_t, A_t) - m(X_t, A_t)) \\ &= \mathbb{1}_{\{E\}} \sum_{t=1}^T \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \mathbb{1}_{\{X_t=x, A_t=a\}} (U_t(x, a) - m(x, a)) \\ &\leq (\sigma_1 + \sigma_2) \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \sum_{t=1}^T \mathbb{1}_{\{X_t=x, A_t=a\}} \sqrt{\frac{8 \log(1/\delta)}{1 \vee N_{x,a}(t-1)}} \\ &\leq (\sigma_1 + \sigma_2) \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \int_0^{N_{x,a}(T)} \sqrt{\frac{8 \log(1/\delta)}{s}} ds \\ &= (\sigma_1 + \sigma_2) \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \sqrt{32N_{x,a}(T) \log(1/\delta)} \leq (\sigma_1 + \sigma_2) \sqrt{32CKT \log(1/\delta)}. \end{aligned}$$

The result follows from choosing $\delta = \frac{1}{T^2}$ and applying the bound $\mathbb{P}[E^c] \leq 2CKT\delta$. \square

E.2 Upper Regret Bounds for poGAMBITTS

The contextual bandit environment of interest with the two-stage reward generating process is fully specified by

$$\mathcal{E}(\theta_1, \theta_2) = \left(P_{\mathcal{X}}, \{\xi_{x,a}^{\theta_1}\}_{x \in \mathcal{X}, a \in \mathcal{A}}, \{m_2(z, x; \theta_2)\}_{x \in \mathcal{X}, a \in \mathcal{A}}, P_\eta \right),$$

where $P_{\mathcal{X}}$ is the context distribution, $\xi_{x,a}^{\theta_1}$ is the distribution of the representation $Z \in \mathbb{R}^d$ given the context-action pair (x, a) , θ_2 is the parameter that governs the mean reward function $m_2(z, x; \theta_2)$, and P_η is the σ_2 -subgaussian noise distribution.

In the analysis of poGAMBITTS, we consider the Bayesian regret with θ_1 fixed to θ_1^* and θ_2 is integrated over the distribution π_2 . We denote the environment with θ_1 fixed to the true parameter θ_1^* by

$$\mathcal{E}(\theta_2) = \mathcal{E}(\theta_1^*, \theta_2) = \left(P_{\mathcal{X}}, \left\{ \xi_{x,a}^{\theta_1^*} \right\}_{x \in \mathcal{X}, a \in \mathcal{A}}, \{m_2(z, x; \theta_2)\}_{x \in \mathcal{X}, a \in \mathcal{A}}, P_\eta \right).$$

Theorem 1 assumes access to an Alg with the Bayesian regret bound of BR_T^{Alg} under the approximate environment

$$\tilde{\mathcal{E}}(\theta_2) = \left(P_{\mathcal{X}}, \{\xi_{x,a}^{off}\}_{x \in \mathcal{X}, a \in \mathcal{A}}, \{m_2(z, x; \theta_2)\}_{x \in \mathcal{X}, a \in \mathcal{A}}, P_\eta \right),$$

where $\xi_{x,a}^{\theta_1^*}$ in the definition of $\mathcal{E}(\theta_2)$ is replaced with its offline empirical estimate $\xi_{x,a}^{off}$ (with density $f_1^{off}(\cdot | x, a)$), and the same mean reward function $m_2(z, x; \theta_2)$ parameterized by θ_2 is used. To be

more precise, the Bayesian regret of the algorithm Alg under the approximate environment $\tilde{\mathcal{E}}(\theta_2)$ is defined as

$$\text{BR}_T^{\text{Alg}} = \mathbb{E}_{\theta_2 \sim \pi_2} \left[\mathbb{E}_{\tilde{\mathcal{E}}(\theta_2)} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t}^{\text{off}}} [m_2(Z, X_t) \mid X_t] - \mathbb{E}_{\xi_{X_t, A_t}^{\text{off}}} [m_2(Z, X_t) \mid X_t] \right) \right] \right],$$

where $\tilde{A}_t^* = \arg \max_a \mathbb{E}_{\xi_{X_t, a}^{\text{off}}} [m_2(Z, X_t) \mid X_t, \theta_2]$; i.e., the optimal action in context X_t under the environment $\tilde{\mathcal{E}}(\theta_2)$. Theorem 1 shows that running the algorithm Alg under the true environment $\mathcal{E}(\theta_2)$ achieves a similar regret bound as long as the approximation of $\mathcal{E}(\theta_2)$ by $\tilde{\mathcal{E}}(\theta_2)$ is sufficiently close.

Theorem 1 (Restated). *Let f_1^{off} be an empirical model for the treatment distribution satisfying (2). Let Alg be a contextual bandit algorithm that selects actions based on the estimated reward function \hat{m}_2 . Then, running Alg in the partially online stochastic treatment setting gives T -step Bayesian regret BR_T^{Alg} , with*

$$\text{BR}_T^{\text{Alg}} \leq \mathcal{O} \left(\widehat{\text{BR}}_T^{\text{Alg}} + T\sqrt{\varepsilon T} \right),$$

where $\widehat{\text{BR}}_T^{\text{Alg}}$ is the T -step Bayesian regret of Alg with respect to \hat{m}_2 .

We provide a proof of Theorem 1 below, which relies on Pinsker's inequality.

Lemma 1 (Pinsker's inequality). *For any pair of measures P and Q on the same probability space (Ω, \mathcal{F}) , we have*

$$\delta(P, Q) = \sup_{A \in \mathcal{F}} P(A) - Q(A) \leq \sqrt{\frac{1}{2} D_{KL}(P \| Q)},$$

where D_{KL} denotes KL divergence in nats.

We now turn to the proof of Theorem 1.

Proof of Theorem 1. Let $P_{x,a}^{\theta_2}$ be the distribution of the reward given the context-action pair (x, a) under the environment $\mathcal{E}(\theta_2)$; i.e., the distribution of the random variable $m_2(Z, x; \theta_2) + \eta$ where $Z \sim \xi_{x,a}^{\theta_1^*}$ and $\eta \sim P_\eta$. Similarly, let $\tilde{P}_{x,a}^{\theta_2}$ be the distribution of the reward given the context-action pair (x, a) under the environment $\tilde{\mathcal{E}}(\theta_2)$; i.e., the distribution of the random variable $m_2(\tilde{Z}, x; \theta_2) + \eta$ where $\tilde{Z} \sim \xi_{x,a}^{\text{off}}$ and $\eta \sim P_\eta$.

Recalling the assumption that f_1^{off} satisfies (2) (i.e., $KL(f_1^{\text{off}}(\cdot; x, a) \| f_1(\cdot; x, a, \theta_1)) \leq \varepsilon$), we observe

$$\begin{aligned} KL(P_{x,a}^{\theta_2} \| \tilde{P}_{x,a}^{\theta_2}) &= KL(\mathcal{L}(g(Z, \eta)) \| \mathcal{L}(g(\tilde{Z}, \eta))) \\ &\leq KL(\mathcal{L}(Z, \eta) \| \mathcal{L}(\tilde{Z}, \eta)) \\ &= KL(\mathcal{L}(Z) \| \mathcal{L}(\tilde{Z})) \\ &\leq \varepsilon, \end{aligned}$$

where $g(z, \eta) = m_2(z, x; \theta_2) + \eta$ and $\mathcal{L}(W)$ denotes the law/distribution of random variable W . The first inequality is due to the data processing inequality.

The interaction between the algorithm Alg and the environment $\mathcal{E}(\theta_2)$ determines the distribution of the sequence $(X_1, A_1, Z_1, Y_1, \dots, X_T, A_T, Z_T, Y_T)$. With a slight abuse of notation, let $P(\theta_2)$ be the probability measure on the sequence generated by Alg under the environment $\mathcal{E}(\theta_2)$ and let $\tilde{P}(\theta_2)$ be the probability measure on the sequence under the environment $\tilde{\mathcal{E}}(\theta_2)$. Then, standard KL divergence decomposition results⁷ for bandit environments gives

$$KL(P(\theta_2) \| \tilde{P}(\theta_2)) = \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} N_{x,a}(T) KL(P_{x,a}^{\theta_2} \| \tilde{P}_{x,a}^{\theta_2}) \leq \varepsilon T,$$

⁷E.g. Lemma 15.1 in Lattimore and Szepesvári [2020] [31].

where $N_{x,a}(T)$ is the count of the pair (x, a) up to time T , as in Appendix E.1. Hence, the Bayesian regret of the algorithm Alg under the environment $\tilde{\mathcal{E}}(\theta_2)$ averaged over the prior distribution on θ_2 is⁸

$$\begin{aligned}
\text{BR}_T &= \mathbb{E}_{\theta_2} \left[\mathbb{E}_{P(\theta_2)} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{\theta_1^*}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{\theta_1^*}} [m_2(Z, X_t; \theta_2)] \right) \right] \right] \\
&\leq \mathbb{E}_{\theta_2} \left[\mathbb{E}_{\tilde{P}(\theta_2)} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{\theta_1^*}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{\theta_1^*}} [m_2(Z, X_t; \theta_2)] \right) \right] \right. \\
&\quad \left. + 2TB \cdot \delta(P(\theta_2), \tilde{P}(\theta_2)) \right] \\
&\leq \mathbb{E}_{\theta_2} \left[\mathbb{E}_{\tilde{P}(\theta_2)} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] \right) \right] \right. \\
&\quad \left. + 2TB \cdot \delta(\xi_{X_t, A_t^*}^{\theta_1^*}, \xi_{X_t, A_t}^{off}) + 2TB \cdot \delta(P(\theta_2), \tilde{P}(\theta_2)) \right] \\
&\leq \mathbb{E}_{\theta_2} \left[\mathbb{E}_{\tilde{P}(\theta_2)} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] \right) \right] \right. \\
&\quad \left. + 2TB \sqrt{\frac{1}{2} KL(\xi_{X_t, A_t^*}^{\theta_1^*} \| \xi_{X_t, A_t}^{off})} + 2TB \sqrt{\frac{1}{2} KL(P(\theta_2) \| \tilde{P}(\theta_2))} \right] \\
&\leq \text{BR}_T^{Alg} + TB \sqrt{2(T+1)\varepsilon}.
\end{aligned}$$

Here, we use the notations $\delta(X, Y)$ for the total variation between X and Y , and $\tilde{A}_t^* = \arg \max_a \mathbb{E}_{\xi_{X_t, a}^{off}} [m_2(Z, X_t; \theta_2) | X_t, \theta_2]$, which is the optimal action in context X_t under the environment $\tilde{\mathcal{E}}(\theta_2)$. The first inequality follows from the fact that the absolute value of the expectant is bounded by $2TB$. The third inequality follows from Pinsker's inequality (Lemma 1). The final inequality follows from the assumption that the algorithm Alg achieves Bayesian regret BR_T^{Alg} under the environment $\tilde{\mathcal{E}}(\theta_2)$, averaged over $\theta_2 \sim \pi_2$. The result follows. \square

E.2.1 Linear Reward Function

Now, we use the reduction result in Theorem 1 to obtain a Bayesian regret bound for poGAMBITS under the environment $\mathcal{E}(\theta_2)$, where the mean reward function $m_2(z, x; \theta_2)$ is linear in z . To apply the reduction, we first establish a Bayesian regret bound for the Thompson sampling algorithm under the environment $\tilde{\mathcal{E}}(\theta_2)$, assuming knowledge of f_1^{off} .

This analysis relies on two standard results from the theory of linear contextual bandits. The first is the concentration bound for the regression-based estimate of the linear parameter:

Lemma 2 (Theorem 20.5 in [31]). *With probability at least $1 - \delta$, we have, for all $t = 1, \dots, T$, that*

$$\|\theta - \hat{\theta}_t\|_{V_t} \leq \beta,$$

where $\beta := 1 + \sigma_2 \sqrt{2 \log(1/\delta) + d \log(1 + T/d)}$.

The second result is a technical lemma commonly used to bound the cumulative sum of upper confidence bounds in the analysis of optimism-based linear contextual bandit algorithms:

Lemma 3 (Elliptical potential lemma). *Let $V_0 \in \mathbb{R}^{d \times d}$ be positive definite and $a_1, \dots, a_n \in \mathbb{R}^d$ be a sequence of vectors with $\|a_t\|_2 \leq L$ for all $t = 1, \dots, n$. Furthermore, let $V_t = V_0 + \sum_{s=1}^t a_s a_s^\top$. Then,*

$$\sum_{t=1}^n \left(1 \wedge \|a_t\|_{V_{t-1}}^2 \right) \leq 2 \log \left(\frac{\det V_n}{\det V_0} \right) \leq 2d \log \left(\frac{\text{Tr}(V_0) + nL^2}{d \det^{1/d}(V_0)} \right).$$

⁸Here, we use $\mathbb{E}_{\theta_2}[W]$ to denote $\mathbb{E}_{\theta_2 \sim \pi_2}[W]$.

Equipped with the two lemmas above, we now derive a Bayesian regret bound for Thompson sampling in the approximate environment $\tilde{\mathcal{E}}(\theta_2)$.

Lemma 4. Consider an environment $\tilde{\mathcal{E}}(\theta_2)$ where the mean reward function is $m_2(z, x; \theta_2) = \langle z, \theta_2 \rangle$. The Thompson sampling algorithm run under this environment with a full knowledge of f_1^{off} , and prior π_2 on θ_2 , achieves a Bayesian regret bound of

$$BR_T \leq \mathcal{O} \left(\sigma_2 d \sqrt{T \log T} \right),$$

Proof. We denote by $\psi_{x,a} = \mathbb{E}_{\xi_{x,a}^{off}} [Z]$. The Bayesian regret can be decomposed as:

$$\begin{aligned} \text{BR}_T &= \mathbb{E} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] \right) \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \langle Z_t^* - Z_t, \theta_2 \rangle \right] + \mathbb{E} \left[\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - Z_t^*, \theta_2 \rangle \right] + \mathbb{E} \left[\sum_{t=1}^T \langle Z_t - \psi_{X_t, A_t}, \theta_2 \rangle \right], \end{aligned}$$

where $A_t^* = \arg \max_a \langle \psi_{X_t, a}, \theta_2 \rangle$ is the best action in context X_t and Z_t^* is a random variable sampled from $\xi_{X_t, A_t^*}^{off}$. The second term above can be bounded by

$$\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - Z_t^*, \theta_2 \rangle \leq \mathcal{O} \left(B \sqrt{T \log \left(\frac{1}{\delta} \right)} \right),$$

which holds with probability at least $1 - \delta$ by the Azuma-Hoeffding inequality, since the sequence $\{\langle \psi_{X_t, A_t^*} - Z_t^*, \theta_2 \rangle\}_{t=1}^T$ forms a martingale difference sequence adapted to the filtration $\mathcal{F}_t = \sigma(\theta_2, X_1, A_1, Z_1, Y_1, \dots, X_t, A_t, Z_t, Y_t, X_{t+1})$. Choosing $\delta = \frac{1}{T}$, we can bound $\mathbb{E} \left[\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - Z_t^*, \theta_2 \rangle \right] \leq \mathcal{O} (B \sqrt{T \log T})$. The third term can be bounded similarly. To bound the first term, we first define an upper confidence bound function $U_t : \mathcal{Z} \rightarrow \mathbb{R}$ as

$$U_t(z) = \langle z, \hat{\theta}_{t-1} \rangle + \beta \|z\|_{V_{t-1}^{-1}},$$

where $V_t = I + \sum_{s=1}^t Z_s Z_s^\top$ and $\hat{\theta}_t = V_t^{-1} \sum_{s=1}^t Y_s Z_s$. By the concentration inequality provided in Lemma 2, we are guaranteed that, for all $t = 1, \dots, T$, $\|\theta_2 - \hat{\theta}_t\|_{V_t} \leq \beta$ with probability at least $1 - \frac{1}{T}$, where we use β defined in the lemma with $\delta = \frac{1}{T}$. Let E_t be the event that $\|\theta_2 - \hat{\theta}_{t-1}\|_{V_{t-1}} \leq \beta$ and $E = \cap_{t=1}^T E_t$. Then,

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T \langle Z_t^* - Z_t, \theta_2 \rangle \right] &= \mathbb{E} \left[\mathbb{1}_{\{E^c\}} \sum_{t=1}^T \langle Z_t^* - Z_t, \theta_2 \rangle \right] + \mathbb{E} \left[\mathbb{1}_{\{E\}} \sum_{t=1}^T \langle Z_t^* - Z_t, \theta_2 \rangle \right] \\ &\leq 2 + \mathbb{E} \left[\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle Z_t^* - Z_t, \theta_2 \rangle \right]. \end{aligned}$$

By the Thompson sampling algorithm, $\mathbb{P}[A_t^* = \cdot | \mathcal{F}_{t-1}] = \mathbb{P}[A_t = \cdot | \mathcal{F}_{t-1}]$, implying $\mathbb{P}[Z_t^* = \cdot | \mathcal{F}_{t-1}] = \mathbb{P}[Z_t = \cdot | \mathcal{F}_{t-1}]$. With this observation, the rest of the proof follows the standard Bayesian regret analysis of linear contextual bandits (e.g., Chapter 36.3 in Lattimore and Szepesvári [2020]) [31].

Since $U_t(z)$ for any fixed $z \in \mathcal{Z}$ is \mathcal{F}_{t-1} -measurable, $\mathbb{E}_{t-1}[U_t(Z_t^*)] = \mathbb{E}_{t-1}[U_t(Z_t)]$.⁹ Therefore, it follows that the second term in the display above is bounded by

$$\begin{aligned} \mathbb{E}_{t-1}[\mathbb{1}_{\{E_t\}} \langle Z_t^* - Z_t, \theta_2 \rangle] &= \mathbb{1}_{\{E_t\}} \mathbb{E}_{t-1}[\langle Z_t^*, \theta_2 \rangle - U_t(Z_t^*) + U_t(Z_t) - \langle Z_t, \theta_2 \rangle] \\ &\leq \mathbb{1}_{\{E_t\}} \mathbb{E}_{t-1}[U_t(Z_t) - \langle Z_t, \theta_2 \rangle] \\ &\leq \mathbb{1}_{\{E_t\}} \langle Z_t, \hat{\theta}_{t-1} - \theta_2 \rangle + \beta \|Z_t\|_{V_{t-1}^{-1}} \\ &\leq 2\beta \|Z_t\|_{V_{t-1}^{-1}}. \end{aligned}$$

Combining the above with $\mathbb{1}_{\{E_t\}} \langle Z_t^* - Z_t, \theta_2 \rangle \leq 2$ produces

$$\begin{aligned} \mathbb{E}\left[\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle Z_t^* - Z_t, \theta_2 \rangle\right] &\leq 2\beta \sum_{t=1}^T (1 \wedge \|Z_t\|_{V_{t-1}^{-1}}) \\ &\leq 2\beta \sqrt{T \mathbb{E}\left[\sum_{t=1}^T (1 \wedge \|Z_t\|_{V_{t-1}^{-1}}^2)\right]} \\ &\leq 2\beta \sqrt{2dT \log(1 + T/d)} \end{aligned}$$

where the second inequality follows from Cauchy-Schwartz and the last inequality from the elliptical potential lemma (Lemma 3). Combining the bounds yields

$$\text{BR}_T \leq \mathcal{O}(\sigma_2 d \sqrt{T \log T}),$$

as required. \square

Corollary 1 in the main body provides a Bayesian regret bound for poGAMBITTS in the linear reward setting. The corollary follows directly by viewing the algorithm as an instance of Thompson sampling for a linear contextual bandit, and applying the reduction in Theorem 1 along with the regret bound in Lemma 4.

E.2.2 General Reward Function

In this section, we analyze the Bayesian regret of poGAMBITTS for general reward functions. We consider function classes \mathcal{F}^{10} that contains the mean reward function $m_2(z, x; \theta_2)$, and derive a regret bound that holds for any such class \mathcal{F} . Our bound is expressed in terms of the eluder dimension of \mathcal{F} , a complexity measure introduced by Russo and Van Roy [2013] [44]. We begin by recalling the definition of the eluder dimension.

Definition 1 (ϵ -dependence). A pair $(z, x) \in \mathcal{Z} \times \mathcal{X}$ is ϵ -dependent on pairs $\{(z_1, x_1), \dots, (z_n, x_n)\} \subseteq \mathcal{Z} \times \mathcal{X}$ with respect to \mathcal{F} if any pair of functions $m, \tilde{m} \in \mathcal{F}$ satisfying $\sqrt{\sum_{i=1}^n (m(z_i, x_i) - \tilde{m}(z_i, x_i))^2} \leq \epsilon$ also satisfies $m(z, x) - \tilde{m}(z, x) < \epsilon$. Further, (z, x) is ϵ -independent of $\{(z_1, x_1), \dots, (z_n, x_n)\}$ with respect to \mathcal{F} if (z, x) is not ϵ -dependent on $\{(z_1, x_1), \dots, (z_n, x_n)\}$.

Definition 2 (Eluder dimension [44]). The ϵ -eluder dimension of a function class \mathcal{F} , $\dim_E(\mathcal{F}, \epsilon)$, is the length d of the longest sequence of pairs in $\mathcal{Z} \times \mathcal{X}$ such that, for some $\epsilon' \geq \epsilon$, every element is ϵ' -independent of its predecessors.

The Bayesian regret of an algorithm under the generator-mediated bandit framework measures the regret of the action A_t taken in context X_t against the best action A_t^* . As a first step toward bounding the Bayesian regret, we bound the regret of the embedding Z_t generated by the action against the counterfactual embedding Z_t^* that would have been generated by the best action. The analysis adapts the proof of Proposition 1 in Russo and Van Roy [2013] ([44]), and expresses the bound in terms of $w_{\mathcal{C}_t}(Z_t, X_t)$ which measures the maximum variation across the reward functions in the confidence set \mathcal{C}_t at (Z_t, X_t) , where

$$w_{\mathcal{C}}(z, x) := \sup_{m, m' \in \mathcal{C}} (m(z, x) - m'(z, x)).$$

⁹Where $\mathbb{E}_t[W] := \mathbb{E}[W | \mathcal{F}_t]$ for any random variable W .

¹⁰For \mathcal{F} consisting of real-valued functions $f : \mathcal{Z} \times \mathcal{X} \rightarrow \mathbb{R}$.

Lemma 5. Fix any sequence of confidence sets $\{\mathcal{C}_t\}_{t=1}^T$ where $\mathcal{C}_t \subseteq \mathcal{F}$ is measurable with respect to $\mathcal{G}_{t-1} = \sigma(X_1, A_1, Z_1, Y_1, \dots, X_{t-1}, A_{t-1}, Z_{t-1}, Y_{t-1}, X_t)$. Then, for any $t = 1, \dots, T$, we have

$$\mathbb{E} \left[\sum_{t=1}^T (m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)) \right] \leq \mathbb{E} \left[\sum_{t=1}^T (w_{\mathcal{C}_t}(Z_t, X_t) + 2B \cdot \mathbb{1}_{\{m_2(\cdot, \cdot; \theta_2) \notin \mathcal{F}\}}) \right].$$

Proof. For any $z \in \mathcal{Z}$ and $x \in \mathcal{X}$, define the upper and lower bounds

$$U_t(z, x) := \sup \{m(z, x) : m \in \mathcal{C}_t\}, L_t(z, x) = \inf \{m(z, x) : m \in \mathcal{C}_t\}.$$

Given any θ_2 , if $m_2(\cdot, \cdot; \theta_2) \in \mathcal{C}_t$, then

$$\begin{aligned} \mathbb{E} [m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)] \\ &\leq \mathbb{E} [U_t(Z_t^*, X_t) - L_t(Z_t^*, X_t) + 2B \cdot \mathbb{1}_{\{m_2(\cdot, \cdot; \theta_2) \notin \mathcal{F}\}}] \\ &\leq \mathbb{E} [w_{\mathcal{C}_t}(Z_t, X_t) + 2B \cdot \mathbb{1}_{\{m_2(\cdot, \cdot; \theta_2) \notin \mathcal{F}\}} + (U_t(Z_t^*, X_t) - U_t(Z_t, X_t))], \end{aligned}$$

By the definition of Thompson sampling, we have $\mathbb{P}[A_t = \cdot | \mathcal{G}_{t-1}] = \mathbb{P}[A_t^* = \cdot | \mathcal{G}_{t-1}]$, implying $\mathbb{P}[Z_t = \cdot | \mathcal{G}_{t-1}] = \mathbb{P}[Z_t^* = \cdot | \mathcal{G}_{t-1}]$.

Since U_t is \mathcal{G}_{t-1} -measurable, $\mathbb{E}[U_t(Z_t^*, X_t) - U_t(Z_t, X_t) | \mathcal{G}_{t-1}] = 0$. It follows that

$$\mathbb{E} \left[\sum_{t=1}^T (m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)) \right] \leq \mathbb{E} \left[\sum_{t=1}^T (w_{\mathcal{C}_t}(Z_t, X_t) + 2B \cdot \mathbb{1}_{\{m_2(\cdot, \cdot; \theta_2) \notin \mathcal{F}\}}) \right],$$

concluding the proof. \square

Russo and Van Roy [2013] ([44]) bounds the sum of the errors $w_{\mathcal{C}_t}(Z_t, X_t)$, where the confidence set \mathcal{C}_t is centered at the least squares estimate \hat{m}_t^{LS} , where

$$\hat{m}_t^{LS} := \arg \min_{m \in \mathcal{F}} \sum_{s=1}^{t-1} (m(Z_s, X_s) - Y_s)^2.$$

The bound, presented in Lemma 6, is in terms of the $\alpha_T^{\mathcal{F}}$ -eluder dimension, where

$$\alpha_t^{\mathcal{F}} := \max \left\{ \frac{1}{t^2}, \inf \{ \|m_1 - m_2\|_{\infty} : m_1, m_2 \in \mathcal{F}, m_1 \neq m_2 \} \right\}.$$

Lemma 6 (Lemma 2 in [44]). If $\{\beta_t\}_{t=1}^T$ is a nondecreasing sequence and $\mathcal{C}_t := \{m \in \mathcal{F} : \|m - \hat{m}_t^{LS}\|_{2, E_t} \leq \sqrt{\beta_t}\}$, then the following holds almost surely:

$$\sum_{t=1}^T w_{\mathcal{C}_t}(Z_t, X_t) \leq \frac{1}{T} + B \cdot \min \{ \dim_E(\mathcal{F}, \alpha_T^{\mathcal{F}}), T \} + 4\sqrt{\dim_E(\mathcal{F}, \alpha_T^{\mathcal{F}})\beta_T T}.$$

Russo and Van Roy [2013] ([44]) also shows that, if the confidence width β_t is chosen as

$$\beta_t^*(\alpha) := 8\sigma_2^2 \log \left(\frac{\mathcal{N}(\mathcal{F}, \alpha, \|\cdot\|_{\infty})}{\delta} \right) + 2\alpha t \left(8B + \sqrt{8\sigma_2^2 \log \left(\frac{4t^2}{\delta} \right)} \right), \quad (3)$$

where $\mathcal{N}(\mathcal{F}, \alpha, \|\cdot\|_{\infty})$ denotes the α -covering number of \mathcal{F} with respect to $\|\cdot\|_{\infty}$, then the associated confidence set (centered at the least squares estimate) contains the true mean reward function $m_2(\cdot, \cdot; \theta_2)$ with high probability. Lemma 7 formalizes this claim.

Lemma 7 (Proposition 2 in [44]). For all $\delta > 0$ and $\alpha > 0$, if

$$\mathcal{C}_t = \left\{ m \in \mathcal{F} : \sum_{s=1}^{t-1} (m(Z_s, X_s) - Y_s)^2 \leq \beta_t^*(\alpha) \right\}$$

for all $t = 1, \dots, T$, then

$$\mathbb{P} \left[m_2(\cdot, \cdot; \theta_2) \in \bigcap_{t=1}^T \mathcal{C}_t \mid \theta_2 \right] \geq 1 - 2\delta.$$

We next establish a Bayesian regret bound for poGAMBITTS in the general reward function setting. As a first step, we analyze the algorithm under the approximate environment $\tilde{\mathcal{E}}(\theta_2)$. We assume that the given function class \mathcal{F} is sufficiently rich to satisfy $\alpha_T^{\mathcal{F}} = T^{-2}$.

Lemma 8. *Consider an environment $\tilde{\mathcal{E}}(\theta_2)$ with mean reward function $m_2(z, x; \theta_2) \in \mathcal{F}$ for some function class \mathcal{F} . The Thompson sampling algorithm run under this environment with a full knowledge of f_1^{off} and a prior π_2 on θ_2 achieves a Bayesian regret bound of*

$$BR_T \leq \mathcal{O} \left(B \cdot \dim_E(\mathcal{F}, T^{-2}) + \sqrt{\dim_E(\mathcal{F}, T^{-2}) \log(\mathcal{N}(\mathcal{F}, T^{-2}, \|\cdot\|_\infty)) T} \right),$$

where the Bayesian regret is defined over the distribution π_2 on θ_2 .

Proof. Well, the Bayesian regret for such a Thompson sampling agent can be decomposed as

$$\begin{aligned} BR_T &= \mathbb{E} \left[\sum_{t=1}^T \left(\mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] \right) \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T (m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)) \right] \\ &\quad + \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - m_2(Z_t^*, X_t; \theta_2) \right] \\ &\quad - \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] - m_2(Z_t, X_t; \theta_2) \right] \\ &\leq \mathbb{E} \left[\sum_{t=1}^T (m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)) \right] + \mathcal{O}(B\sqrt{T \log T}), \end{aligned}$$

where the final inequality follows from the Azuma-Hoeffding inequality on the martingale difference sequences

$$\left\{ \mathbb{E}_{\xi_{X_t, A_t^*}^{off}} [m_2(Z, X_t; \theta_2)] - m_2(Z_t^*, X_t; \theta_2) \right\}_{t=1}^T$$

and

$$\left\{ \mathbb{E}_{\xi_{X_t, A_t}^{off}} [m_2(Z, X_t; \theta_2)] - m_2(Z_t, X_t; \theta_2) \right\}_{t=1}^T,$$

both of which are adapted to $\mathcal{G}_t = \sigma(X_1, A_1, Z_1, Y_1, \dots, X_t, A_t, Z_t, Y_t, X_{t+1})$.

As before, we consider $A_t^* := \arg \max_a \mathbb{E}_{\xi_{X_t, a}^{off}} [m_2(Z, X_t; \theta_2) \mid X_t, \theta_2]$ and we consider Z_t^* sampled from $\xi_{X_t, A_t^*}^{off}$. Considering a fixed sequence of confidence sets $\mathcal{C}_1, \dots, \mathcal{C}_T$, as defined in Lemma 6, and applying Lemma 5 gives the following bound on the first term:

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T (m_2(Z_t^*, X_t; \theta_2) - m_2(Z_t, X_t; \theta_2)) \right] &\leq \mathbb{E} \left[\sum_{t=1}^T (w_{\mathcal{C}_t}(Z_t, X_t) + 2B \cdot \mathbb{1}_{\{m_2(\cdot, \cdot; \theta_2) \notin \mathcal{F}\}}) \right] \\ &\leq \mathcal{O} \left(B \cdot \dim_E(\mathcal{F}, \alpha_T^{\mathcal{F}}) + \sqrt{\dim_E(\mathcal{F}, \alpha_T^{\mathcal{F}}) \beta_T^*(\alpha_T^{\mathcal{F}}) T} \right), \end{aligned}$$

where the second inequality follows by Lemma 6 and Lemma 7, with $\delta = \frac{1}{T}$. The desired bound follows by plugging $\alpha = \alpha_T^{\mathcal{F}}$ into the definition of $\beta_T^*(\alpha)$ (Equation 3). \square

Finally, Theorem 2 directly follows by the reduction provided in Theorem 1 and the lemma above.

Theorem 2 (Restated). *Let f_1^{off} satisfy (2). For any function class \mathcal{F} with $\hat{m}_2 \in \mathcal{F}$, the Bayesian regret of poGAMBITTS (Algorithm 2) satisfies*

$$BR_T^{po:\mathcal{F}} \leq \tilde{\mathcal{O}} \left(\sigma_2 \sqrt{\dim_E(\mathcal{F}, T^{-2}) \log \mathcal{N}(\mathcal{F}, T^{-2}, \|\cdot\|_\infty) T} + T\sqrt{\varepsilon T} \right).$$

E.2.3 Bounding the Error in Offline Treatment Models

poGAMBITTS requires access to an estimate $f_1^{off}(\cdot; a, x)$ that is ε -close to the true treatment density $f_1(\cdot; a, x, \theta_1)$ in KL divergence. Below, we consider the example where $\xi_{x,a}^{\theta_1}$ is Gaussian for each $(x, a) \in \mathcal{X} \times \mathcal{A}$, demonstrating that a sample-based empirical estimate can meet this requirement.

Lemma 9 (Multivariate Gaussian Distribution). *Consider a multivariate normal distribution $\mathcal{N}(\mu, \Sigma)$ and let $\{Z_1, \dots, Z_n\}$ be i.i.d. samples drawn from this distribution. Then, the estimated distribution $\mathcal{N}(\hat{\mu}, \hat{\Sigma})$ satisfies*

$$KL\left(\mathcal{N}(\mu, \Sigma) \parallel \mathcal{N}(\hat{\mu}, \hat{\Sigma})\right) \leq \mathcal{O}\left(\frac{d}{\lambda_{\min}(\Sigma)} \sqrt{\frac{d + \log(1/\delta)}{n}}\right)$$

with probability at least $1 - \delta$, where $\hat{\mu}$ is the sample mean and $\hat{\Sigma}$ is the empirical covariance matrix.

Proof. The KL divergence has the following exact form:

$$KL\left(\mathcal{N}(\mu, \Sigma) \parallel \mathcal{N}(\hat{\mu}, \hat{\Sigma})\right) = \frac{1}{2} \left[\text{tr}(\hat{\Sigma}^{-1} \Sigma) + (\hat{\mu} - \mu)^\top \hat{\Sigma}^{-1} (\hat{\mu} - \mu) - d + \log \det \hat{\Sigma} - \log \det \Sigma \right].$$

We now turn to bounding each term. A standard concentration result for empirical covariance matrices, based on the matrix Bernstein inequality, gives:

$$\|\hat{\Sigma} - \Sigma\| \leq \sqrt{\frac{d + \log(1/\delta)}{n}}$$

with probability at least $1 - \delta$, where $\|\cdot\|$ is the operator norm. Hence, the inverse $\hat{\Sigma}$ concentrates and it follows that

$$\|\hat{\Sigma}^{-1}\| \leq 2\|\Sigma^{-1}\| = \frac{2}{\lambda_{\min}(\Sigma)}$$

for $n \geq \Omega\left(\frac{d + \log(1/\delta)}{\lambda_{\min}(\Sigma)^2}\right)$.

Hence, we can bound

$$\begin{aligned} \text{tr}(\hat{\Sigma}^{-1} \Sigma) &= \text{tr}(I + \hat{\Sigma}^{-1}(\Sigma - \hat{\Sigma})) \\ &\leq d + \|\hat{\Sigma}^{-1}\| \|\Sigma - \hat{\Sigma}\| \\ &\leq d + \mathcal{O}\left(\frac{1}{\lambda_{\min}(\Sigma)} \sqrt{\frac{d + \log(1/\delta)}{n}}\right). \end{aligned}$$

To bound the log-determinant term, observe that

$$\begin{aligned} \log \det \hat{\Sigma} - \log \det \Sigma &= \log \det(I + \Sigma^{-1/2}(\hat{\Sigma} - \Sigma)\Sigma^{-1/2}) \\ &\leq d\|\Sigma^{-1/2}(\hat{\Sigma} - \Sigma)\Sigma^{-1/2}\| \\ &\leq d\|\Sigma^{-1}\| \|\hat{\Sigma} - \Sigma\| \\ &\leq \mathcal{O}\left(\frac{d}{\lambda_{\min}(\Sigma)} \sqrt{\frac{d + \log(1/\delta)}{n}}\right). \end{aligned}$$

Finally, we bound

$$\begin{aligned} (\hat{\mu} - \mu)^\top \hat{\Sigma}^{-1} (\hat{\mu} - \mu) &\leq \|\hat{\mu} - \mu\|_2^2 \|\hat{\Sigma}^{-1}\| \\ &\leq \mathcal{O}\left(\frac{d + \log(1/\delta)}{n}\right), \end{aligned}$$

applying the vector Bernstein inequality to bound $\|\hat{\mu} - \mu\|_2$.

Combining all the bounds yields

$$KL\left(\mathcal{N}(\mu, \Sigma) \parallel \mathcal{N}(\hat{\mu}, \hat{\Sigma})\right) \leq \mathcal{O}\left(\frac{d}{\lambda_{\min}(\Sigma)} \sqrt{\frac{d + \log(1/\delta)}{n}}\right),$$

as required. \square

E.3 Upper Regret Bound for foGAMBITTS

In this section, we derive a Bayesian regret for foGAMBITTS in the environment $\mathcal{E}(\theta_1, \theta_2)$. The key step in the analysis is that the action selected by Thompson sampling has the same conditional distribution as the optimal action, given the agent's posterior. Formally, let $\mathcal{F}_t = \sigma(X_1, A_1, Z_1, Y_1, \dots, X_t, A_t, Z_t, Y_t, X_{t+1})$ be the σ -algebra generated by the interaction up to time t . Thompson sampling logic guarantees that the conditional distributions of A_t^* and A_t given \mathcal{F}_{t-1} are identical; i.e.,

$$\mathbb{P}[A_t^* = \cdot | \mathcal{F}_{t-1}] = \mathbb{P}[A_t = \cdot | \mathcal{F}_{t-1}], \quad \text{almost surely.} \quad (4)$$

We will use the following concentration inequality to control deviations of a martingale sequence in the proof of Theorem 3:

Lemma 10 (Freedman's inequality). *Let (\mathcal{F}_k) be a filtration and let the sequence $\{X_k\}$ of random variables satisfy $\mathbb{E}[X_k | \mathcal{F}_{k-1}] = 0$ and $|X_k| \leq R$ almost surely. Let $S_k = \sum_{i=1}^k X_i$ and $V_k = \sum_{i=1}^k \mathbb{E}[X_i^2 | \mathcal{F}_{i-1}]$. Then, with probability at least $1 - \delta$, we have, for all k , that*

$$|S_k| \leq \sqrt{2V_k \log(2/\delta)} + \frac{2R}{3} \log(2/\delta).$$

We now restate Theorem 3 and proceed with its proof.

Theorem 3 (Restated). *If $m_2(z, x_t; \theta)$ is linear in $z \in \mathbb{R}^d$, the Bayesian regret of foGAMBITTS (Algorithm 1) satisfies*

$$BR_T^{fo:lin} \leq \tilde{\mathcal{O}}\left(\sigma_1 \sqrt{CKT} + \sigma_2 d \sqrt{T}\right).$$

Proof of Theorem 3. We are interested in bounding the Bayesian regret defined as

$$BR_T = \mathbb{E}\left[\sum_{t=1}^T \mathbb{E}_{\xi_{X_t, A_t^*}^{\theta_1}} [m_2(Z_t, X_t; \theta_2)] - \mathbb{E}_{\xi_{X_t, A_t}^{\theta_1}} [m_2(Z_t, X_t; \theta_2)]\right]$$

where $A_t^* = \arg \max_{a \in \mathcal{A}} \mathbb{E}_{\xi_{X_t, a}^{\theta_1}} [m_2(Z, X_t; \theta_2) | X_t, \theta_2]$ is the best action in context X_t . Under the linear reward model $m_2(z, x; \theta_2) = \langle z, \theta_2 \rangle$, we can write the Bayesian regret as

$$BR_T = \mathbb{E}\left[\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right]$$

where we define $\psi_{x,a} := \mathbb{E}_{\xi_{x,a}^{\theta_1}}[Z]$, suppressing the superscript θ_1 for convenience.

Since the agent observes the embedding Z_t and receives a noisy reward of the form $Y_t = \langle Z_t, \theta_2 \rangle + \eta_t$, it can estimate θ_2 by regressing Y_t on Z_t . Consider the following estimate for θ_2 :

$$\hat{\theta}_t = V_t^{-1} \sum_{s=1}^t Y_s Z_s$$

where $V_t = I + \sum_{s=1}^t Z_s Z_s^\top$. Lemma 2 guarantees $\|\theta_2 - \hat{\theta}_t\|_{V_t} \leq \beta$ for all $t = 1, \dots, T$ with probability at least $1 - \delta$, where β , dependent on δ , is defined in the lemma. Let E_t be the event that $\|\theta_2 - \hat{\theta}_{t-1}\|_{V_{t-1}} \leq \beta$, and let $E := \cap_{t=1}^T E_t$. Setting $\delta = \frac{1}{T}$ in the definition of β gives $\mathbb{P}[E] \geq 1 - \frac{1}{T}$. We can then bound the Bayesian regret under E^c and E separately, as follows.

$$\begin{aligned} BR_T &= \mathbb{E}\left[\sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right] \\ &= \mathbb{E}\left[\mathbb{1}_{\{E^c\}} \sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right] + \mathbb{E}\left[\mathbb{1}_{\{E\}} \sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right] \\ &\leq 2 + \mathbb{E}\left[\mathbb{1}_{\{E\}} \sum_{t=1}^T \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right] \\ &\leq 2 + \mathbb{E}\left[\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle\right]. \end{aligned} \quad (5)$$

Before proceeding with the bound, we clarify the probability space governing the random variables X_t, A_t, Z_t, Y_t . We adopt the random table model (see Chapter 4.6 in Lattimore and Szepesvári [2020]) and, conditional on θ_1 , define the embedding $Z_t(x, a) \sim \xi_{x,a}^{\theta_1}$ independently for each $t = 1, \dots, T$ and $(x, a) \in \mathcal{X} \times \mathcal{A}$ [31]. We define $Z_t = Z_t(X_t, A_t)$ to be the embedding observed by the agent. This random table model provides a convenient framework for the analysis that follows.

To proceed, we define the upper confidence bound function $U_t : \mathcal{Z} \rightarrow \mathbb{R}$ by

$$U_t(z) = \langle z, \widehat{\theta}_{t-1} \rangle + \beta \|z\|_{V_{t-1}^{-1}}.$$

Then, under the event E_t , it follows that, for all $(x, a) \in \mathcal{X} \times \mathcal{A}$,

$$\begin{aligned} \langle \psi_{x,a}, \theta_2 \rangle - U_t(Z_t(x, a)) &= \langle \psi_{x,a}, \theta_2 \rangle - \langle Z_t(x, a), \widehat{\theta}_{t-1} \rangle - \beta \|Z_t(x, a)\|_{V_{t-1}^{-1}} \\ &= \langle \psi_{x,a} - Z_t(x, a), \theta_2 \rangle + \langle Z_t(x, a), \theta_2 - \widehat{\theta}_{t-1} \rangle - \beta \|Z_t(x, a)\|_{V_{t-1}^{-1}} \\ &\leq \langle \psi_{x,a} - Z_t(x, a), \theta_2 \rangle. \end{aligned}$$

The inequality follows from using Cauchy-Schwarz to bound $\langle Z_t(x, a), \theta_2 - \widehat{\theta}_{t-1} \rangle \leq \|Z_t(x, a)\|_{V_{t-1}^{-1}} \|\theta_2 - \widehat{\theta}_{t-1}\|_{V_{t-1}}$. Furthermore, under the event E_t , $\|\theta_2 - \widehat{\theta}_{t-1}\|_{V_{t-1}} \leq \beta$. Similarly,

$$\begin{aligned} U_t(Z_t(x, a)) - \langle \psi_{x,a}, \theta_2 \rangle &= \langle Z_t(x, a), \widehat{\theta}_{t-1} \rangle + \beta \|Z_t(x, a)\|_{V_{t-1}^{-1}} - \langle \psi_{x,a}, \theta_2 \rangle \\ &= \langle Z_t(x, a), \widehat{\theta}_{t-1} - \theta_2 \rangle + \beta \|Z_t(x, a)\|_{V_{t-1}^{-1}} - \langle \psi_{x,a} - Z_t(x, a), \theta_2 \rangle \\ &\leq \langle Z_t(x, a) - \psi_{x,a}, \theta_2 \rangle + 2\beta \|Z_t(x, a)\|_{V_{t-1}^{-1}}. \end{aligned}$$

Since $U_t(\cdot)$ and X_t are \mathcal{F}_{t-1} -measurable, (4) implies

$$\mathbb{P}[U_t(Z_t(X_t, A_t^*)) = \cdot \mid \mathcal{F}_{t-1}] = \mathbb{P}[U_t(Z_t(X_t, A_t)) = \cdot \mid \mathcal{F}_{t-1}],$$

and that

$$\mathbb{E}_{t-1}[U_t(Z_t(X_t, A_t^*))] = \mathbb{E}_{t-1}[U_t(Z_t(X_t, A_t))],$$

where, as before, $\mathbb{E}_{t-1}[\cdot]$ denotes $\mathbb{E}[\cdot \mid \mathcal{F}_{t-1}]$. Hence, recalling the definition $Z_t = Z_t(X_t, A_t)$,

$$\begin{aligned} \mathbb{E}_{t-1}[\mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle] &= \mathbb{1}_{\{E_t\}} \mathbb{E}_{t-1} [\langle \psi_{X_t, A_t^*}, \theta_2 \rangle - U_t(Z_t(X_t, A_t^*)) + U_t(Z_t(X_t, A_t)) - \langle \psi_{X_t, A_t}, \theta_2 \rangle] \\ &\leq \mathbb{1}_{\{E_t\}} \mathbb{E}_{t-1} [\langle \psi_{X_t, A_t^*} - Z_t(X_t, A_t^*), \theta_2 \rangle + \langle Z_t - \psi_{X_t, A_t}, \theta_2 \rangle + 2\beta \|Z_t\|_{V_{t-1}^{-1}}]. \end{aligned}$$

Continuing the regret bound from (5) and using $\mathbb{E}_{t-1}[\mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle] \leq 2$, we observe

$$\begin{aligned} \text{BR}_T &\leq 2 + \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{t-1} [\mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - \psi_{X_t, A_t}, \theta_2 \rangle] \right] \\ &\leq 2 + \mathbb{E} \left[\underbrace{\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - Z_t(X_t, A_t^*), \theta_2 \rangle}_{(a)} \right] \\ &\quad + \mathbb{E} \left[\underbrace{\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle Z_t - \psi_{X_t, A_t}, \theta_2 \rangle}_{(b)} \right] + 2\beta \mathbb{E} \left[\underbrace{\sum_{t=1}^T (\|Z_t\|_{V_{t-1}^{-1}} \wedge 3)}_{(c)} \right] \end{aligned}$$

To bound term (a), we express it as a summation over (x, a) :

$$\begin{aligned} \sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - Z_t(X_t, A_t^*), \theta_2 \rangle &= \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \sum_{t=1}^T \mathbb{1}_{\{E_t\}} \mathbb{1}_{\{X_t=x, A_t^*=a\}} \langle \psi_{x,a} - Z_t(x, a), \theta_2 \rangle. \end{aligned}$$

For any given context-action pair $(x, a) \in \mathcal{X} \times \mathcal{A}$, the sequence $\{W_t(x, a)\}_{t=1}^T$, defined by

$$W_t(x, a) = \mathbb{1}_{\{E_t\}} \mathbb{1}_{\{X_t=x, A_t^*=a\}} \langle Z_t(x, a) - \psi_{x,a}, \theta_2 \rangle$$

is a martingale difference sequence adapted to

$$\mathcal{G}_t = \sigma(\theta_1, \theta_2, X_1, A_1, \{Z_1(x, a)\}_{x,a}, Y_1, \dots, X_t, A_t, \{Z_t(x, a)\}_{x,a}, Y_t, X_{t+1}, A_{t+1}).$$

This follows because $\mathbb{1}_{\{E_t\}}$ and $\mathbb{1}_{\{X_t=x, A_t^*=a\}}$ are \mathcal{G}_{t-1} -measurable, implying

$$\begin{aligned} & \mathbb{E} [\mathbb{1}_{\{E_t\}} \mathbb{1}_{\{X_t=x, A_t^*=a\}} \langle Z_t(x, a) - \psi_{x,a}, \theta_2 \rangle \mid \mathcal{G}_{t-1}] \\ &= \mathbb{1}_{\{E_t\}} \mathbb{1}_{\{X_t=x, A_t^*=a\}} \langle \mathbb{E}[Z_t(x, a) - \psi_{x,a} \mid \mathcal{G}_{t-1}], \theta_2 \rangle \\ &= 0. \end{aligned}$$

Furthermore, $|W_t(x, a)| \leq 2$ for all $t = 1, \dots, T$, and

$$\begin{aligned} \sum_{t=1}^T \mathbb{E}[W_t(x, a)^2 \mid \mathcal{G}_{t-1}] &\leq \sum_{t=1}^T \mathbb{1}_{\{X_t=x, A_t^*=a\}} \mathbb{E}[(\langle Z_t(x, a) - \psi_{x,a}, \theta_2 \rangle)^2 \mid \mathcal{G}_{t-1}] \\ &\leq \sigma_1^2 \sum_{t=1}^T \mathbb{1}_{\{X_t=x, A_t^*=a\}} \\ &= \sigma_1^2 N_T(x, a), \end{aligned}$$

where $N_T(x, a) = \sum_{t=1}^T \mathbb{1}_{\{X_t=x, A_t^*=a\}}$, as before. The first inequality follows from the assumption that $m_2(Z, x; \theta_2)$ is σ_1 -subgaussian. Hence, Freedman's inequality (Lemma 10) coupled with a union bound over $(x, a) \in \mathcal{X} \times \mathcal{A}$, guarantees that, with probability at least $1 - \delta$,

$$\sum_{t=1}^T W_t(x, a) \leq \sigma_1 \sqrt{2N_T(x, a) \log\left(\frac{2CK}{\delta}\right)} + \frac{4}{3} \log\left(\frac{2CK}{\delta}\right).$$

It follows that

$$\begin{aligned} \sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t^*} - Z_t(X_t, A_t^*), \theta_2 \rangle &= \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \sum_{t=1}^T W_t(x, a) \\ &\leq \sum_{x \in \mathcal{X}} \sum_{a \in \mathcal{A}} \left(\sigma_1 \sqrt{2N_T(x, a) \log\left(\frac{2CK}{\delta}\right)} + \frac{4}{3} \log\left(\frac{2CK}{\delta}\right) \right) \\ &\leq \sigma_1 \sqrt{2CKT \log\left(\frac{2CK}{\delta}\right)} + \frac{4}{3} CK \log\left(\frac{2CK}{\delta}\right). \end{aligned}$$

where the last inequality is by Cauchy-Schwarz and the fact that $\sum_x \sum_a N_T(x, a) = T$. Similar analysis, with A_t in place of A_t^* , yields the following bound on the term (b):

$$\sum_{t=1}^T \mathbb{1}_{\{E_t\}} \langle \psi_{X_t, A_t} - Z_t(X_t, A_t), \theta \rangle \leq \sigma_1 \sqrt{2CKT \log\left(\frac{2CK}{\delta}\right)} + \frac{4}{3} CK \log\left(\frac{2CK}{\delta}\right).$$

Finally, the last term (c) can be bounded using the elliptical potential lemma (Lemma 3) as

$$\begin{aligned} \sum_{t=1}^T (3 \wedge \|Z_t\|_{V_{t-1}^{-1}}) &\leq 3 \sum_{t=1}^T (1 \wedge \|Z_t\|_{V_{t-1}^{-1}}) \\ &\leq 3 \sqrt{T \sum_{t=1}^T (1 \wedge \|Z_t\|_{V_{t-1}^{-1}}^2)} \\ &\leq 3 \sqrt{2dT \log\left(1 + \frac{T}{d}\right)}. \end{aligned}$$

Combining the bounds above, we observe

$$\begin{aligned} \text{BR}_T &\leq 2 + \sigma_1 \sqrt{8CKT \log\left(\frac{2CK}{\delta}\right)} + \frac{8}{3}CK \log\left(\frac{2CK}{\delta}\right) + 6\beta \sqrt{2dT \log\left(1 + \frac{T}{d}\right)} \\ &\leq \mathcal{O}\left(\sigma_1 \sqrt{CKT \log\left(\frac{CA}{\delta}\right)} + \sigma_2 d \sqrt{T \log\left(1 + \frac{T}{d}\right) \left(\log\left(1 + \frac{T}{d}\right) + \log\left(\frac{1}{\delta}\right)\right)} + CK \log\left(\frac{CA}{\delta}\right)\right). \end{aligned}$$

The result follows. □

E.4 Lower Regret Bound for foGAMBITTS

In this section, we establish a minimax lower bound on the regret of foGAMBITTS.

Claim 2. *The minimax regret of foGAMBITTS satisfies*

$$R_T \geq \Omega\left(\max\left\{\sigma_1 \sqrt{CKT}, \sigma_2 \sqrt{dT \log K}\right\}\right).$$

Proof. To establish the lower bound, we construct instances in which the agent cannot efficiently resolve uncertainty, leading to provably high regret. We first show $R_T \geq \Omega\left(\sigma_1 \sqrt{CKT}\right)$.

Fix K, C and d , and let $\mathcal{A} = \{1, \dots, K\}$, $\mathcal{X} = \{1, \dots, C\}$. Let $\theta_2 = (\sigma_1, 0, \dots, 0) \in \mathbb{R}^d$ be the parameter for the linear reward model. Furthermore, choose the reward noise $\sigma_2 = 0$.

Let $\theta_1 \in \{\mathcal{A} \cup \{0\}\}^C$ parameterize the distributions $\xi_{x,a}^{\theta_1}$ such that, for each $(x, a) \in \mathcal{X} \times \mathcal{A}$,

$$\xi_{x,a}^{\theta_1} \sim \text{Bernoulli on } \{e_1, \mathbf{0}\} \text{ with probabilities } \begin{cases} \left(\frac{1}{2}, \frac{1}{2}\right) & \text{if } a \neq \theta_{1,x}, \\ \left(\frac{1}{2} - \Delta, \frac{1}{2} + \Delta\right) & \text{if } a = \theta_{1,x}, \end{cases}$$

for some $\Delta \in [0, \frac{1}{4}]$.¹¹ Here, we use e_1 to denote the first basis vector in \mathbb{R}^d .

Under this environment, the mean reward of each context-action pair is σ_1 -subgaussian. Also, in context x , the mean reward of every action is $\frac{\sigma_1}{2}$ except for one optimal action $a = \theta_{1,x}$, which gives mean reward $\frac{1}{2}(\sigma_1 + \Delta)$.

Because rewards in different contexts are generated from independent distributions, observations collected in one context carry no statistical information about any other context. Hence, from the learner's point of view, the problem decomposes into C independent instances of standard K -armed bandits, each with horizon $N_T(1), \dots, N_T(C)$, where $N_T(x)$ denotes the number of rounds in which context x appears.

For a context x , the regret lower bound over $N_T(x)$ rounds is $c\sigma_1 \sqrt{KN_T(x)}$ for some constant c (see Chapter 15.2 in Lattimore and Szepesvári [2020] [31]). Summing over x , and choosing $N_T(x) = \frac{T}{C}$ to ensure the environment chooses contexts evenly, we obtain

$$R_T \geq \Omega\left(\sum_{x \in \mathcal{X}} \sqrt{KN_T(x)}\right) = \Omega\left(\sigma_1 \sqrt{CKT}\right).$$

We now show that $R_T \geq \Omega\left(\sigma_2 \sqrt{dT \log K}\right)$. Consider an environment where $\xi_{x,a}^{\theta_1}$ is deterministic and known to the learner. This reduces to a standard contextual linear bandit environment with dimension d . Applying a known lower bound for this setting [32], we obtain

$$R_T \geq \Omega\left(\sigma_2 \sqrt{dT \log K}\right).$$

¹¹Note that $\{\mathcal{A} \cup \{0\}\}^C$ here denotes the C -dimensional space with components elements of $\{\mathcal{A} \cup \{0\}\}$, rather than the complement of $\{\mathcal{A} \cup \{0\}\}$.

Combining the two lower bounds yields

$$R_T \geq \Omega \left(\max \left\{ \sigma_1 \sqrt{CKT}, \sigma_2 \sqrt{dT \log K} \right\} \right),$$

as required. \square

F Further Simulations

This appendix presents additional simulations designed to probe specific factors influencing GAMBITTS performance. Unless otherwise specified, all data-generating environments assume a linear outcome model with optimism, formality, and severity as predictors.

F.1 Variance Decomposition

Section 5 and Appendix E discuss GAMBITTS performance in terms of the relative variance of the treatment and outcome processes. In our setup, the treatment is generated by an LLM and fixed in advance, but we can control the residual variance of the outcome-generative model. This allows us to manipulate the ratio of treatment variance to outcome noise directly. Throughout the main simulations, we held this ratio fixed by setting $\sigma_1 = \sigma_2$.¹² In Figure 7, we increase σ_2 to explore how performance changes under high outcome noise. We scale σ_2 gradually, up to the level observed when fitting the outcome-generative model on held-out data from the 2023 IHS study.

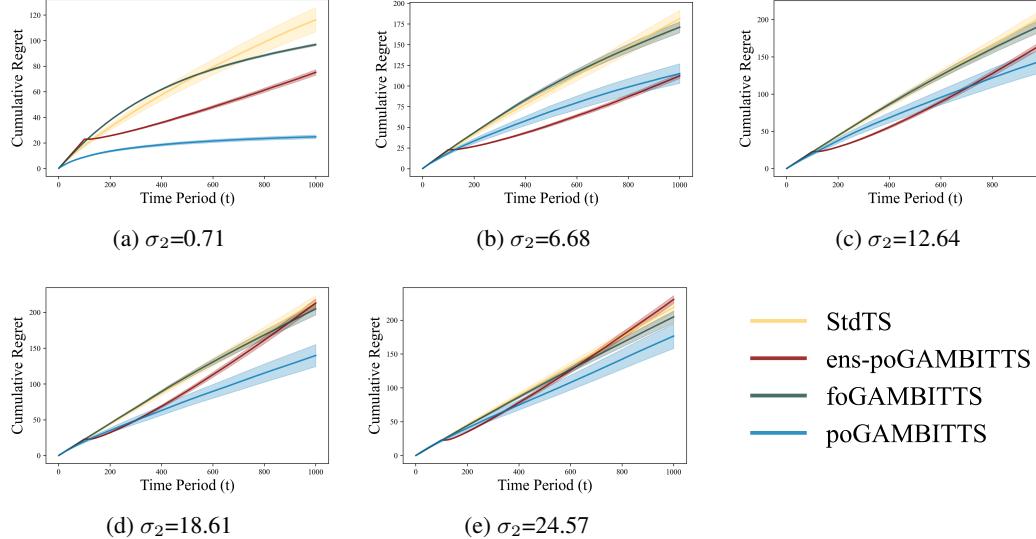


Figure 7: Cumulative Regret Across Outcome Variances

Corollary 1 and Theorem 3 suggest that the performance advantage of linear GAMBITTS-based algorithms over standard Thompson sampling increases with the outcome noise level σ_2 . Figure 7 shows that linear GAMBITTS-based methods achieve lower regret across values of σ_2 ; however, as σ_2 increases, all algorithms exhibit performance deterioration and begin to converge. One possible explanation is that higher outcome noise reduces the overall signal-to-noise ratio, making it more difficult for *any* bandit agent to quickly learn a good policy. Future work may examine this phenomenon over longer time horizons to better understand the asymptotic behavior of the competing methods.

¹²We computed σ_1 from the generated text interventions used for outcome simulation; i.e., from the response_db described in Appendix G.

F.2 Nonlinear Data-Generative Reward Model

Section 6.2 examined GAMBITTS performance under style embedding misspecification. Figure 8 explores a related question, focusing instead on misspecification in the structure of the outcome model. In these experiments, all GAMBITTS agents used the correct style embeddings, but the true outcome-generating process $\mathbb{E}[Y | Z, X]$ followed an MLP structure. To assess how GAMBITTS agents respond to this nonlinear specification, we revisit the variance-scaling analysis from Figure 7, now under an MLP outcome model.

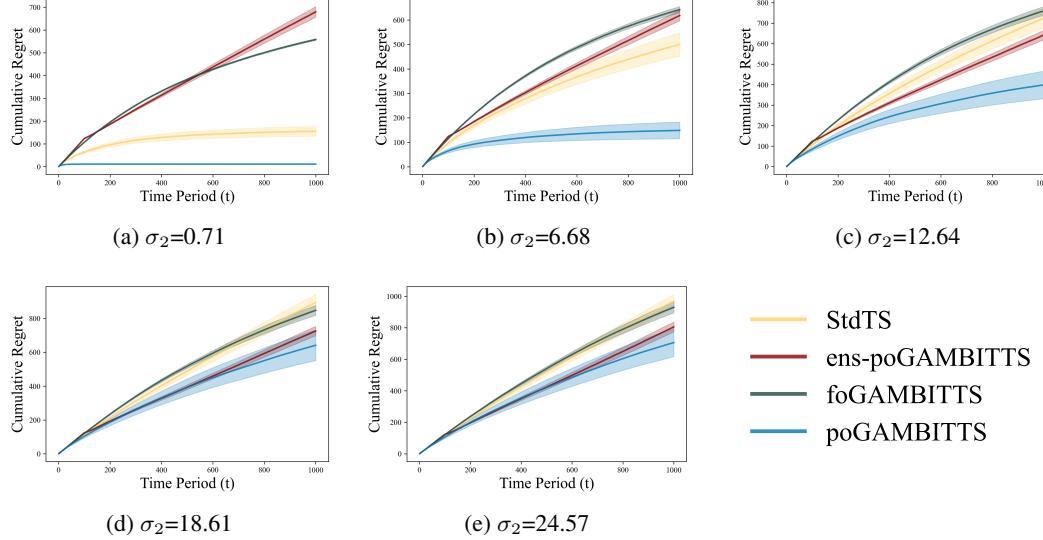


Figure 8: Cumulative Regret Across Outcome Variances: MLP Outcome Generative Model

In a nonlinear data-generating environment, we expected `ens-poGAMBITTS` to outperform its linear counterparts, given its more flexible outcome model. However, as in other simulation settings, `ens-poGAMBITTS` exhibits volatile performance. Interestingly, the linear `poGAMBITTS` agent performs relatively well, despite the misspecification of the outcome model. One possible explanation for this pattern, and for the broader volatility observed in `ens-poGAMBITTS` performance, is that, within the GAMBITTS framework (and in Thompson sampling more generally), the outcome model is used to support decision-making rather than to produce accurate predictions. As described in Appendix G.2, the neural network models are substantially more parameterized than their linear counterparts, and may require far more data to converge. In contrast, while the linear models are misspecified, they may still capture sufficient structure to support effective decision-making.

The volatility in `ens-poGAMBITTS` performance suggests that the approach holds promise, but it is not a panacea. Without tuning tailored to the specifics of the application, the agent may perform poorly. Future work could develop principled strategies for hyperparameter selection to make `ens-poGAMBITTS` more robust to variation in the data-generating process.

E.3 poGAMBITTS Simulator Access

As discussed in Section 4, partially online GAMBITTS variants use simulator access to construct f_1^{off} , an empirical estimate of the treatment-generating distribution. Figure 9 examines the sensitivity of `poGAMBITTS` and `ens-poGAMBITTS` performance to the number of samples per (x, a) pair used to construct f_1^{off} .

We expected the partially online GAMBITTS variants to perform better with increased access to the treatment-generating distribution. Figure 9 supports this hypothesis: cumulative regret decreases as the number of simulator draws increases. Interestingly, performance stabilizes after roughly 50 draws, suggesting that limited simulation access may be sufficient to realize most of the benefits of the partially online approach.

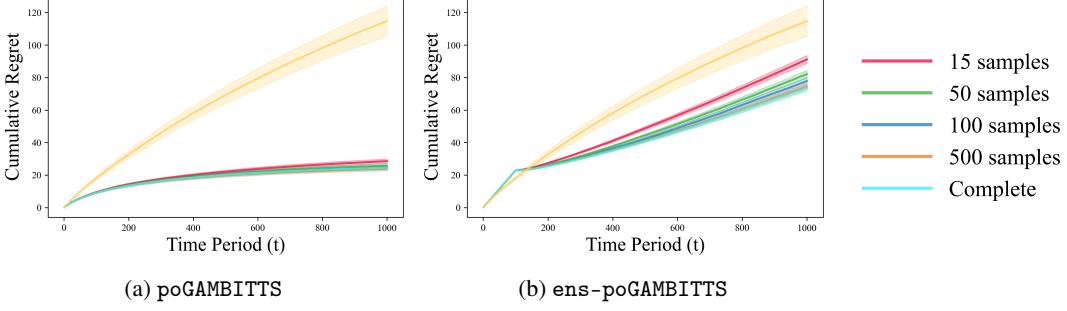


Figure 9: Cumulative Regret Under Varying Levels of Simulator Access

F.4 Treatment Dimension

We also examine the effect of d , the dimensionality of the treatment embedding. To do so, we introduce 15 additional style dimensions, supplementing the five used in Section 6, for a total of 20 dimensions: optimism, formality, encouragement, severity, clarity, humor, complexity, vision, detail, threat-level, urgency, politeness, personalization, conciseness, actionability, emotiveness, authoritativeness, authenticity, supportiveness, and gender-codedness.¹³

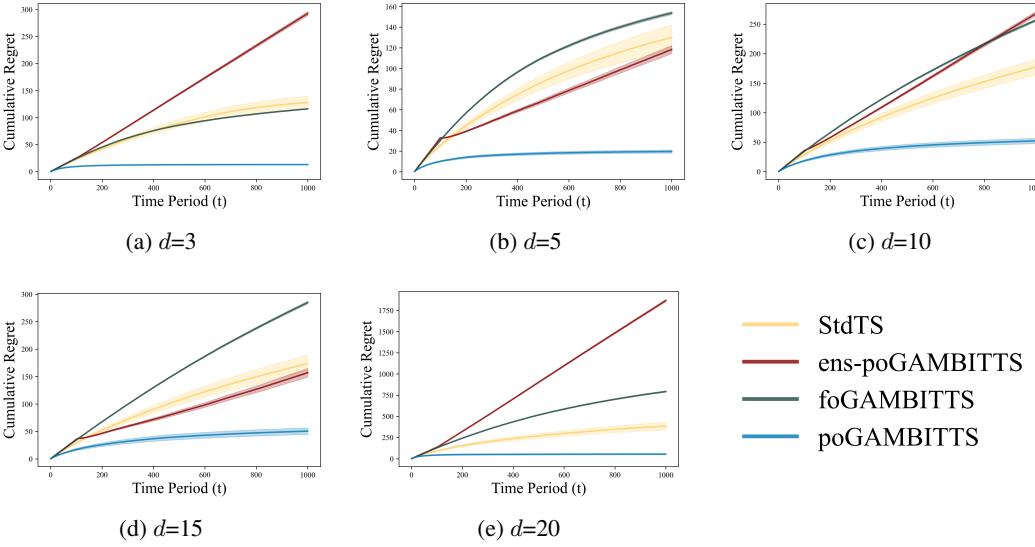


Figure 10: Cumulative Regret for Varying Treatment Embedding Dimension

We anticipated that GAMBITTS agents would perform increasingly poorly relative to standard Thompson sampling as the embedding dimension d increased. Figure 10 confirms this trend for the foGAMBITTS agent. However, the linear poGAMBITTS agent remains competitive even as d grows. As discussed in Appendix F.2, the volatility in ens-poGAMBITTS performance further underscores the need for future work on robust hyperparameter tuning strategies within the GAMBITTS framework.

F.5 Direct Contextual Influence on Reward

The simulations in Sections 6.1 and 6.2 (along with Appendix C) use reward models linear in Z^{optimism} with Gaussian noise. As described in Section 6.3, the other data-generating environments

¹³To select the dimensions used in Figures 10a-10e, we clustered the style dimensions by their correlation in `response_db`. For each d , we used d -means clustering and randomly sampled one dimension from each cluster.

extend this structure to include Z^{optimism} , $Z^{\text{formality}}$, and Z^{severity} , again with Gaussian noise.¹⁴ This appendix considers a simulation designed to evaluate GAMBITTS performance when rewards depend directly on context as well as semantic dimensions.¹⁵ Given the substantial performance advantage of GAMBITTS agents over standard Thompson sampling in our main experiments, we expect this advantage to persist under additional covariate influence.

In this appendix, we extend the setup from Figure 4b¹⁶ by allowing rewards to depend directly on context. As described in Appendix G, the simulated contextual variables were current location and steps taken the previous day. Specifically, we use the data-generating model¹⁷

$$Y_t = \beta^{\text{optimism}} Z_t^{\text{optimism}} + \beta^{\text{formality}} Z_t^{\text{formality}} + \beta^{\text{severity}} Z_t^{\text{severity}} \\ + \beta^{\text{steps}} X_t^{\text{stepsprevday}} + \sum_{\text{loc} \in \text{Locs}} \beta^{\text{loc}} \mathbb{1}_{X_t^{\text{currloc}} = \text{loc}} + \varepsilon_t.$$

Figure 11 reports results under correctly specified GAMBITTS agents. The figure supports the intuition that the advantage of GAMBITTS persists under direct covariate influence, showing little deviation from Figure 4b.¹⁸

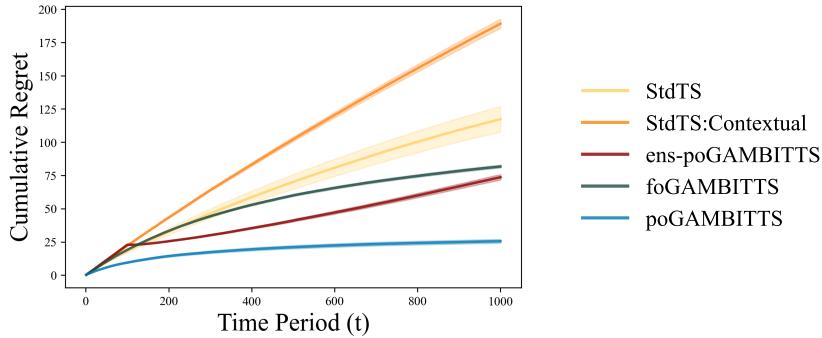


Figure 11: Cumulative Regret Under Reward Model with Direct Context Influence

F.6 Scaling Number of Arms: Revisited

As observed in Section 6.3, foGAMBITTS underperforms standard Thompson sampling when the number of available arms, K , is large. However, the more appropriate comparison is to a contextual Thompson sampler, since foGAMBITTS incorporates covariate information through the treatment embedding. Moreover, Section 5 shows that foGAMBITTS is expected to outperform standard Thompson sampling when $\sigma_2 \gg \sigma_1$, whereas the experiments in Section 6.3 assume $\sigma_1 = \sigma_2$. To investigate this further, we examine foGAMBITTS performance with $K = 40$ under varying levels of outcome noise.

¹⁴With two exceptions: Appendix F.2 instead uses an MLP mean reward model, and Appendix F.4 varies the number of semantic dimensions in the mean reward model.

¹⁵In all cases, rewards depend on text indirectly through the semantic dimensions. Semantic dimensions are functions of LLM-produced text. Although other data-generative environments do not specify a direct effect of context on reward, context enters implicitly through the LLM query, which shapes the generated text and hence the semantic dimensions (described further in Appendix G).

¹⁶We use Figure 4b as the baseline because it employs the default choices for (i) the number of arms, (ii) the set of style dimensions, and (iii) the reward noise level used throughout the experiments in this manuscript.

¹⁷Here $\text{Locs} := \{\text{"home"}, \text{"work"}, \text{"a location other than home or work"}\}$ and $X^{\text{stepsprevday}}$ was recoded to be 1, 2, 3, or 4 if the simulated previous day's steps was 0-4,999, 5,000-9,999, 10,000-15,000, or more than 15,000, respectively.

¹⁸The standard contextual Thompson sampler continues to underperform in this setting. Incorporating covariates requires a more heavily parameterized model for the contextual Thompson sampling agent than for its non-contextual counterpart. When covariate effects are modest, this added complexity increases variance and can reduce performance (a known issue in contextual bandit learning [47]). Furthermore, in this simulation, the optimal action does not vary across context, limiting the potential gain from contextual information. Further work could examine GAMBITTS performance under different forms and strengths of contextual influence.

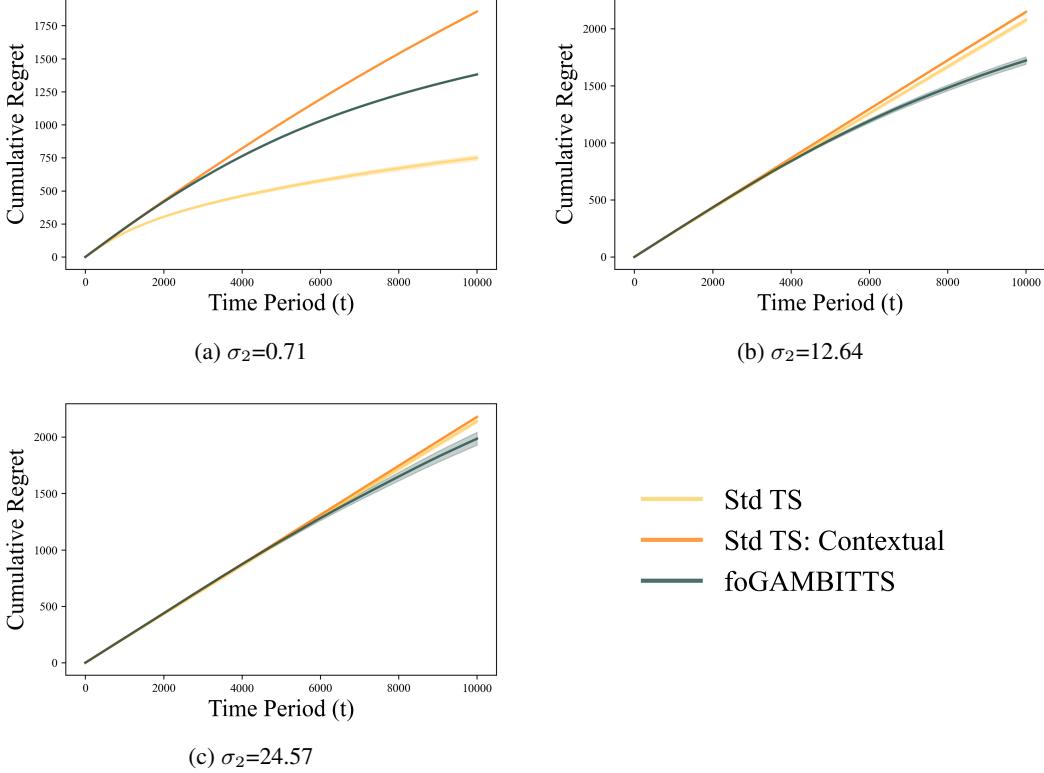


Figure 12: Cumulative Regret for foGAMBITTS with $K = 40$ under Varying Reward Noise

Figure 12 supports the results in Section 5: foGAMBITTS performs increasingly well relative to standard Thompson sampling as σ_2 grows. Moreover, even for small σ_2 , foGAMBITTS outperforms a contextual Thompson sampling agent. To better observe asymptotic behavior under high noise, we extend the time horizon to 10,000 rounds. Even so, cumulative regret does not appear to level off, suggesting that the agents are still learning the optimal policy. This persistent learning phase likely reflects the low signal-to-noise ratio estimated from the IHS study that generated the outcome parameters. Future research directions to address such scenarios could incorporate partial pooling methods (e.g., Huch et al. [2020]) to share reward model information across users and accelerate policy learning [20].

G Simulation Study Design

To evaluate algorithm performance in a setting where ground truth is available, we needed full knowledge of the data-generating process. This ruled out querying a black-box language model during simulation runs, since doing so would obscure the underlying conditional expectations and make performance comparisons difficult to interpret. Instead, we simulated on-the-fly text generation by constructing a response database (referred to as `response_db`): for each prompt-context pair, we sampled 1,000 responses from an LLM¹⁹ and computed the style embeddings associated with each response, as detailed in Appendix G.3. At runtime, we drew from this pool to mimic LLM-generated responses, as described in Algorithm 6.

¹⁹In particular, Llama 3.1 with 8.0B parameters, 131,072 context length, 2,095 embedding length, and Q4_K_M quantization, accessed via Ollama [42]. Additionally, we re-ran select analyses using a Qwen-based LLM. The results were consistent with those in the main manuscript, suggesting that the simulation conclusions are not sensitive to the choice of LLM, though further work would be needed to establish this more formally.

We derived the outcome-generating model using data from the 2023 IHS study.²⁰ For each outcome model structure (linear or neural network) and relevant style dimensions, we fit the corresponding model to the IHS data and used the estimated parameters to define the outcome surface. This provided a data-driven but controlled setting for comparing policies.²¹ We provide the exact parameters used in the data-generating process on GitHub.

Lastly, in our simulations, we included two context variables:

$$\begin{aligned} \text{stepsprevday} &\in \{"0\text{-}4,999", "5,000\text{-}9,999", "10,000\text{-}15,000", "more than 15,000"\} \text{ and} \\ \text{currloc} &\in \{"\text{home}", "\text{work}", "\text{a location other than home or work}"\}. \end{aligned}$$

These context variables varied uniformly across time points and were sampled independently of the action.

Algorithm 6 Pseudocode for Outcome Generation

- Inputs:** `response_db`, style embedding functions f_1, \dots, f_D , mean model structure $m_2(z, x; \theta_2)$, data-generative outcome model parameters θ_2^{IHS} , conditional variance σ_2^2 .
- 1: **Given** a prompt a and context x , randomly sample response g from `response_db`, restricted to responses generated from (a, x) .
 - 2: **Compute** style embedding vector $z = [f_1(g) \dots f_D(g)]^\top$.
 - 3: **Simulate** outcome y by drawing $\varepsilon \sim \mathcal{N}(0, \sigma_2)$ and setting $y = m_2(z, x; \theta_2^{IHS}) + \varepsilon$.
-

G.1 Prompt Specifications

To initialize the LLM environment for crafting intervention text, we provided the following system prompt:

You are tasked with writing a single message to an individual. The message should be two to three sentences long and will be delivered to the individual through a mobile app. The intent of the message is to encourage the individual to take more steps, with the ultimate goal of improving their physical health. The app collects the following information on its users: whether they took 0-4,999, 5,000-9,999, 10,000-15,000, or more than 15,000 steps the previous day, and whether they are currently at home, at work, or at another location.

To generate the intervention text used in our simulations, we constructed prompts that explicitly manipulated specific stylistic dimensions. Each prompt was designed to target a particular user context and asked the LLM to produce a message in which a given dimension (e.g., optimism, formality) was either high or low. This approach yielded a total of 40 prompts (one high and one low version for each of the 20 dimensions) attempting to promote alignment among the variation in the generated text with the interpretable axes used in the outcome model.

To generate simulation messages (for given contexts `stepsprevday` and `currloc`), we used the following prompt structure:

The app provides the following context about the individual: they took `stepsprevday` steps the previous day and are currently at `currloc`. This context may be included in the message but does not need to dominate it. Please make the message mildly `<style polarity>. <dimension description>`. Please do not include anything in your response other than the text of the message.

Here, `<style polarity>` refers to the high/low level of the target stylistic dimension (e.g., “optimistic”/“pessimistic” for the optimism dimension). `<dimension description>` is a descriptor

²⁰We used the square root of the number of steps taken by the user on the day following the message as the reward. This transformation, rather than using raw step counts, yielded approximately normally distributed errors.

²¹The one exception is our use of $\sigma_1 = \sigma_2$ in the main simulations. We observed σ_1 from the variability in `response_db` and explored sensitivity to larger values of σ_2 in Appendix F.1.

of the stylistic quality, intended to guide the LLM toward the appropriate tone.²² Unless otherwise noted (e.g., Section 6.3), all experiments were run using five prompts available to the agent, chosen to be evenly distributed by expected outcome.²³

G.2 Agent Specifications

We evaluate several bandit agents in Section 6 and Appendix F. Unless otherwise noted, each agent type followed the base specifications described in the corresponding subsection below. Furthermore, all GAMBITTS-based algorithms used a correctly specified mean outcome model (unless otherwise noted; e.g., Section 6.2). When possible, we used data from the 2022 IHS study to inform prior specification. For example, the average square root of daily steps was approximately 77, which guided our choice of prior for intercept terms.

G.2.1 Standard Thompson Samplers

The standard Thompson sampling agents modeled the reward for each action using a normal likelihood, with a prior of the form $\mu_a \sim \mathcal{N}(77, \sigma_a^2)$, where $\sigma_a^2 \sim \text{Inverse-Gamma}(1, 10)$. The contextual Thompson sampling agents used the same model and priors, but indexed them by each context-action pair (x, a) .

G.2.2 poGAMBITTS Agent

Linear poGAMBITTS agents modeled the reward as $Y \sim \mathcal{N}(\beta_0 + Z^\top \beta', \Sigma)$ for $Z, \beta' \in \mathbb{R}^d$ and $\Sigma \in \mathbb{R}^{d \times d}$. Letting $\beta := [\beta_0 \ \beta']^\top \in \mathbb{R}^{d+1}$, these agents placed the following priors on model parameters:

$$\begin{aligned}\theta_2 &\sim \mathcal{N}(\mu_0, B_0) \\ \mu_0 &= [77 \ 0 \dots 0]^\top \in \mathbb{R}^{d+1} \\ B_0 &= \text{diag}(10^{-2}, 1, \dots, 1) \in \mathbb{R}^{(d+1) \times (d+1)}.\end{aligned}$$

G.2.3 ens-poGAMBITTS Agent

Our implementation of ens-poGAMBITTS uses PyTorch [1], with each neural network instantiated using PyTorch’s default initialization: weights and biases were drawn from a uniform distribution $\mathcal{U}\left(-\frac{1}{\sqrt{k}}, \frac{1}{\sqrt{k}}\right)$, where k is the number of input features. The ensembles consisted of $M_{ens} = 60$ networks. Each network was single-layer feedforward model with 64 hidden units and ReLU activation. Online training was performed with a batch size of 100 and a learning rate of 0.1. The ensembles maintained a replay buffer of size 1,024. To approximate the expectation in Algorithm 4, we used 100 Monte Carlo samples. In each experiment, ens-poGAMBITTS training began after a burn-in period of $t = 100$ steps, during which the neural networks did not update, allowing sufficient data to accumulate for batch-based optimization. Lastly, the agents used the data-generative σ_2 as the perturbation noise standard deviation.

G.2.4 foGAMBITTS Agent

foGAMBITTS agents used the same prior π_2 on θ_2 as defined for the poGAMBITTS agent in Appendix G.2.2 above. foGAMBITTS agents used a normal likelihood for $Z \in \mathbb{R}^d$ with mean $\theta_1 \in \mathbb{R}^d$. The prior π_1 for θ_1 is given below:

$$\begin{aligned}\Sigma_1 &\sim \text{Inverse-Wishart}(\mathbf{I}_d, 1) \\ \theta_1 &\sim \mathcal{N}(\mathbf{0}_d, \Sigma_1),\end{aligned}$$

where $\mathbf{I}_d \in \mathbb{R}^{d \times d}$, $\mathbf{0}_d \in \mathbb{R}^d$, and d is the dimension of the latent treatment space \mathcal{Z} .

²²For example, the following is the description of the optimism dimension: “Here, optimism is defined as the degree of hopeful or confident language in the message. A lower optimism level may include a more matter-of-fact or pessimistic tone, while a higher optimism level may include uplifting and buoyant language.”

²³E.g., if an experiment included only three prompts, we selected the highest, lowest, and median prompt based on true expected value under the data-generating process.

G.3 Text Embedding Construction

We aimed to quantify LLM-generated responses along five stylistic dimensions: optimism, encouragement, formality, clarity, and severity.²⁴ For each dimension, d , we constructed a mapping $f_d : \mathcal{G} \rightarrow \mathbb{R}$, as described in Algorithm 7 below. This appendix details the construction of these embeddings and provides evidence that they capture the intended stylistic dimensions. This evidence is included primarily to aid interpretability of the simulation experiments. While high-quality embeddings are crucial in applied settings, they are not required for demonstrating the validity of the GAMBITTS approaches.

We learned each f_d by prompting an LLM to generate a set of JITAI-style messages (brief behavioral messages aimed at encouraging users to walk more, mirroring the setup of the IHS and HeartSteps studies) [33, 40]. Appendix G.3.1 below contains the exact prompt template used for this purpose. For each prompt, we asked the LLM to vary the degree to which its response reflected a particular stylistic dimension (d). The responses were intended to isolate variation along that dimension while holding other properties relatively constant. The goal was not to assign explicit scores or labels, but to induce embeddings in which each axis captured targeted semantic variation.

We then trained a variational autoencoder (VAE) on the resulting text set, constraining the model to use a single latent dimension. The output of this latent space was used as the embedding $f_d(g)$ for each message g .²⁵ This process was repeated separately for each dimension. Algorithm 7 below provides pseudocode for this process.

Algorithm 7 Pseudocode for Style Dimension d Embedding Construction

- Input:** Style dimension label d , LLM, high-dimensional embedding model, VAE architecture, training size M_{train}^d .
- 1: **Generate** M_{train}^d random prompts $\{p_i^d\}_{i=1}^{M_{train}^d}$ using the template presented in Appendix G.3.1.
 - 2: **For** each prompt, p_i^d , in Step 1, query the LLM with the prompt to receive r_i^d
 - 3: **For** each response, r_i^d in Step 2, obtain a high-dimensional numerical embedding $(e_i^{d,h})$ by using the high-dimensional embedding model [55].
 - 4: **Given** embeddings $e_1^{d,h}, \dots, e_{M_{train}^d}^{d,h}$, train a variational auto-encoder to obtain a one-dimensional latent representation of the response [6]. Call the resulting mapping f_d .
-

As with the outcome generative model, we used Llama 3.1 as the LLM for embedding construction. Additionally, we used BERT as our high-dimensional embedding model [10], and set $M_{train}^d = 22,000$.

G.3.1 Embedding Generation Prompt Creation

To initialize the LLM environment, we provided a system prompt describing the context and purpose of the message generation task (e.g., encouraging physical activity through behavioral messages). This global instruction was held fixed across all prompt constructions and consisted of the following:

You are tasked with writing eleven separate messages to an individual. Each message should be two to three sentences long and will be delivered to the individual through a mobile app. The messages should be independent and must not reference each other. The intent of each message is to encourage the individual to take more steps, with the ultimate goal of improving their physical health. The app collects the following information on its users: whether they took 0-4,999, 5,000-9,999, 10,000-15,000, or more than 15,000 steps the previous day, and whether they are currently at home, at work, or at another location.

To simulate messages that varied along a given stylistic axis, we used a template-based user prompt, shown below. The specific wording of the prompt varied by stylistic dimension; we present the version used for *optimism* as an illustrative example.

²⁴Along with the additional 15 dimensions for the experiments in Appendix F.4.

²⁵We note that the messages used to train the VAE were distinct from those contained in the `response_db` database used for outcome simulation.

The app provides the following context about the individual: they took `stepsprevday` steps the previous day and are currently at `currloc`. This context may be included in the messages but should not dominate them.

The primary axis of variation in the messages should be optimism, with no intentional variation along other dimensions such as tone, length, or formality. For this task, create eleven messages where the optimism level varies as follows: Message 1 should represent the least optimistic tone (optimism level 0). Message 11 should represent the most optimistic tone (optimism level 10). Messages in between should gradually and evenly increase in optimism. Here, optimism is defined as the degree of hopeful or confident language in the message. A lower optimism level may include a more matter-of-fact or pessimistic tone, while a higher optimism level may include uplifting and buoyant language. Remember, low levels of optimism should have a pessimistic tone.

Please write each message on a separate line, and remember that the target length is two to three sentences for each message. Do not include anything in your response other than the text of the messages. Start every new message with its number. The optimism of the message should be independent of the user specific information.

For this purpose, we randomly drew context variables `stepsprevday` and `currloc` uniformly from their possible values.²⁶ We varied the user context across prompts to encourage the VAE to learn stylistic variation independently of context. For each style dimension d , we generated 2,000 responses (each consisting of eleven JITAI messages spanning the stylistic axis) yielding a total of 22,000 messages per dimension.

G.3.2 Text Embedding Results

After running Algorithm 7, we compared the target input rating with the VAE output. Figure 13 shows this relationship.²⁷

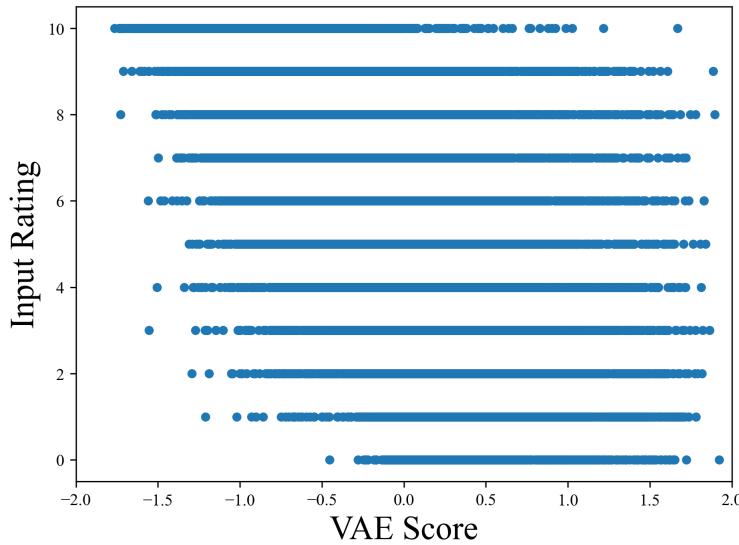


Figure 13: Optimism Scores versus Prompted Rating

We find that the VAE captured some signal related to the intended tone of the message. While the relationship in Figure 13 is noisy, two caveats are worth noting. (i) Precisely quantifying variation along a semantic dimension is likely a domain-specific challenge and, while important, it is outside the scope of this work. Our goal here is not to validate the embedding as a standalone

²⁶We used the same context variables for the embedding construction as for the outcome generation.

²⁷Note that we removed two outliers with high VAE scores, one with rating ten and one with rating nine, which distorted the scale of the plot.

construct; nonetheless, alignment between prompted tone and latent score adds interpretability to the simulation setup. (ii) Some of the observed noise reflects inconsistencies in LLM behavior: the model occasionally generates responses that do not align with the requested tone. For example, the following message was produced in response to a prompt with optimism rating 1, but appears highly optimistic: “You’re on a roll after yesterday’s success - keep up the pace and see where it takes you!” Larger LLMs or more carefully tuned unsupervised methods may improve the correspondence between prompted tone and latent score.

G.3.3 Embedding Construction Summary

Appendix G.3.1 presents the prompt instructions used for generating text that varied in optimism, and Appendix G.3.2 shows how the resulting VAE scores relate to the specified input levels. As illustrated in Figure 13, the VAE captures variation along the target semantic dimension, but its unsupervised nature means the sign of the scores is arbitrary (for example, positive scores appear to align with pessimistic rather than optimistic text). Table 1 summarizes the style definitions provided in the prompts for each of the five primary semantic dimensions and reports the styles of text the VAE associated with positive and negative scores.²⁸

Table 1: Primary Semantic Styles, Prompted Definitions, and Interpretation of VAE Scores

Style	Style Definition	VAE Score Interpretations	
		Negative	Positive
Optimism	“The degree of hopeful or confident language in the message. A lower optimism level may include a more matter-of-fact or pessimistic tone, while a higher optimism level may include uplifting and buoyant language.”	Optimistic	Pessimistic
Severity	“The degree of dire or drastic language in the message. A lower severity level may include a more lax, calm, or gentle tone, while a higher severity level may include more dark, intense, worrying and distressing language.”	Severe	Lax
Formality	“The degree of which the message has an objective, academic, or professional tone. A lower formality level may include a more personal, casual, or emotional tone and include colloquial language. A higher formality level may include more matter-of-fact, impersonal, professional and serious language.”	Informal	Formal
Clarity	“The degree to which the intent of the message is intelligibly communicated. A lower clarity level should be more vague and ambiguous in its language, while a higher clarity level may include more precise, coherent, and intelligible instructions or comments.”	Clear	Unclear
Encouragement	“The degree of persuasion using positivity, confidence, and hope. A lower encouragement level may include a more depressed, dispirited or hopeless tone, while a higher encouragement level may include more excited, heartening, and motivating language.”	Discouraging	Encouraging

G.3.4 Relationship between Style Embeddings

Table G.3.4 below shows the empirical correlation (within the `response_db` dataset constructed for outcome generation) among the five stylistic embeddings discussed in Section 6.2.

²⁸These associations were identified by manually examining a sample of generated text and its corresponding VAE scores.

Table 2: Style Dimension Correlation Table

	Optimism	Formality	Severity	Clarity	Encouragement
Optimism	1.000	-0.095	0.787	0.184	-0.960
Formality	-0.095	1.000	-0.568	-0.248	-0.121
Severity	0.787	-0.568	1.000	0.066	-0.644
Clarity	0.184	-0.248	0.066	1.000	-0.201
Encouragement	-0.960	-0.121	-0.644	-0.201	1.000

Furthermore, we analyzed the structure of the relationship between the different style dimensions. Figure 14 shows pairwise scatterplots of style embeddings for the five primary dimensions. Recall that, in addition to the five primary dimensions, we created 15 more for the experiments in Appendix F.4. The pairwise scatterplots for all 20 dimensions are shown in Figure 15. These figures provide a qualitative view of the dependencies among dimensions. However, the direction of the embeddings is not directly comparable across dimensions; e.g., a high encouragement score corresponds to an encouraging tone, while a high optimism score reflects greater pessimism.

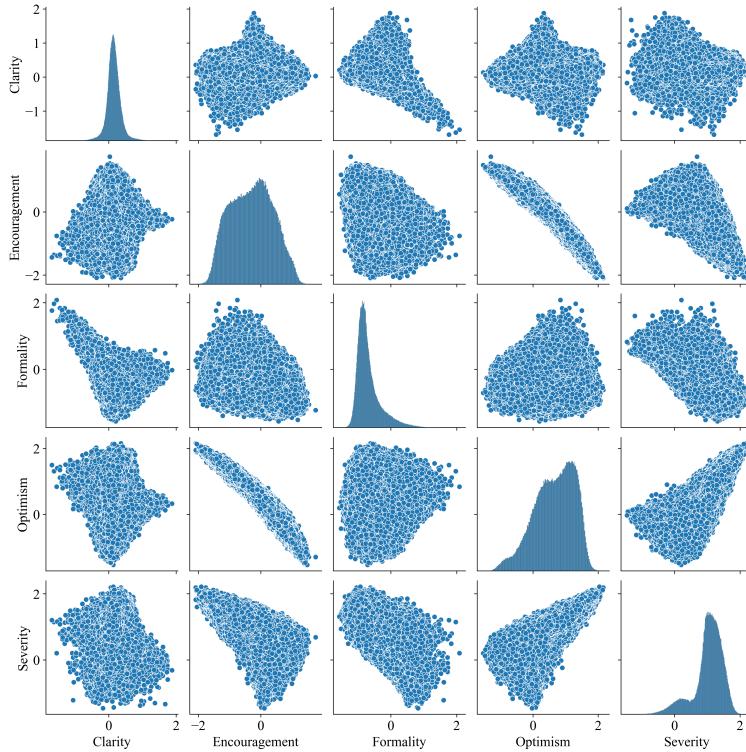


Figure 14: Pairwise Scatterplot of Five Primary Style Dimensions

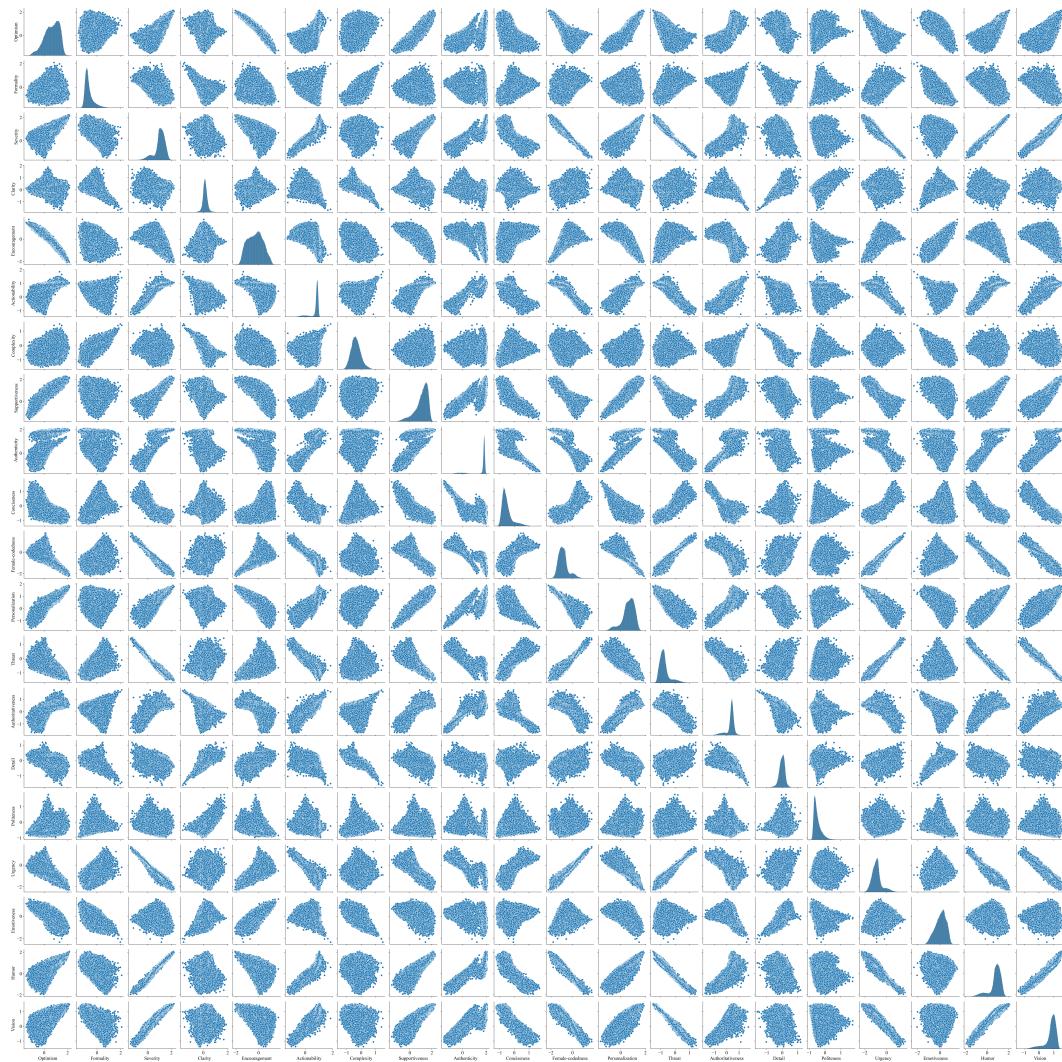


Figure 15: Pairwise Scatterplot of All Style Dimensions

G.4 Computing Resources

The simulation of text-based treatments generated by Llama 3.1 was implemented on a high-performance compute cluster. Each node in the cluster included two 2.9 GHz Intel Xeon Gold 6226R processors, 8 GB of allocated RAM, and a single NVIDIA A40 GPU with 48 GB of memory. VAE training was performed on a similar setup, except each node was allocated 16 GB of RAM. The simulation studies were conducted on nodes equipped with two 3.0 GHz Intel Xeon Gold 6154 processors, 16 CPU cores, and between 16–32 GB of allocated RAM. Under these specifications, each experiment presented in Section 6 and Appendix F required approximately 50–70 minutes to complete.

H Broader Impact

As discussed in Section 1, text-based sequential interventions are used across domains such as mobile health and education to support behavioral change and improve outcomes. We view the GAMBIT framework and associated algorithms as a step toward more personalized decision-making in such settings. At the same time, personalization introduces risks (particularly in sensitive domains) where automatically generated content may influence individuals in subtle or unintended ways. In some applications, it may be inappropriate to deliver GenAI-generated content without human oversight or safety review. We discuss these concerns and potential adaptations to the framework in Section 7, with the hope that variants of GAMBIT can be developed to address such challenges directly.